

Index

1. Introduction	3
1.1 Pre-requisites	4
2. Infrastructure basics.....	4
2.1 Create a Security Group	4
2.1.1 Create the initial inbound rules	4
2.1.2 Update the Security Group	5
2.1.3 What the hell did we just do	5
2.1.4 Important note regarding SSH inbound:	5
2.1.5 Standard vs custom ports.....	5
2.1.6 TCP vs UDP protocols	6
2.1.7 SG-Summary.....	6
2.2 Launch EC2 Instances	6
2.2.1 First let us create the instances.....	6
2.2.1.1 ..and this script does what?	7
2.2.2 Summary, launch and multiples.....	7
2.2.3 Note Instance IPs.....	7
2.2.4 EC2-Summary.....	7
2.3 Initialize Docker Swarm via SSH	8
2.3.1 Create a nifty helper-document (optional but recommended)	8
2.3.2 Connect to the Manger and initialize Swarm.....	8
2.3.2.1 I got most of that. I think	8
2.3.3 Let us add the Worker nodes	9
2.3.3.1 Let us quickly verify the Swarm cluster.....	9
2.3.4 Swarm-Summary.....	9
2.4 Deploying the services manually	10
2.4.1 Create a Docker Compose file	10
2.4.1.1 Let us not go into all details, but roughly what this does is this	10
2.4.2 Deploy the freshly created Stack	11
2.4.3 Verify the deployment	11
2.4.4 Deploy-Summary.....	11
2.5 Test web service and see if the service can scale.....	12
2.5.1 Let us first verify web access.....	12
2.5.2 Scale the services	13

2.5.3 Tools to monitor the services	14
2.5.4 Web-Summary	14
2.6 Tools for cleaning up (optional).....	14
2.6.1 Remove stack	14
2.6.2 Leave Swarm	14
2.6.3 Terminate Instances	14
2.6.4 Delete Security Group	15
2.7 Summary of above sections	15
3 CloudFormation me up, Scotty	15
3.1 S3 Bucket roles first.....	16
3.1.1 Create the bucket first.....	16
3.1.2 Then attach the policy to the user so you can use the bucket.....	16
3.2 Quick conversion – CloudFormation style.....	17
3.2.1 Nested stacks structure.....	17
3.2.2 What why when how all those cookies, conceptually? IaC is cool	17
3.2.3 root.yaml (Master stack)	18
3.2.4 00-sg-swarm.yaml (Swarm Security Group).....	20
3.2.5 10-ec2-swarm.yaml (EC2 Manager + 2 Workers, Docker installed)	21
3.2.6 Nested stacks – big picture (TL;DR).....	23
3.3 Important note on SSH, dynamic IP ranges, ingress/egress	24
3.3.1 This is important to know re: SSH / AllowedSshCidr.....	24
3.3.2 Let us also mention egress rules	24
3.4 CF-Summary	25
4. Deploying our CF and validating it	25
4.1 Tiny quality of life life-hacks.....	25
4.1.1 Optional, tiny .env for copy-pasting shorter CLI.....	25
4.1.2 Create dev.params.json	25
4.1.2.1 Important note, do not forget to set values	26
4.1.2.2 Where to find the VPC ID	26
4.1.2.3 Where to find the Subnet ID and which to pick and why	26
4.1.2.4 Using CLI to solve the above	27
4.2 Validate template (syntax check)	27
4.3 Upload the templates to S3	28
4.4 Create the stack (daily bring-up of the stacks).....	28
4.4.1 Grab outputs (IPs, SG id)	28

4.4.2 Error handling when creating.....	29
4.5 Update the stack (apply template changes)	29
4.6 Tear-down of the stacks (daily or frequent routine)	30
4.7 Optional: Check the public IPs of the Instances.....	30
4.7.1 Bonus: Capture the three public IPs into variables	31
4.7.2 Temporary: Until we automate deploy, do it manually.....	31
4.7.3 Deploying the Docker Compose (temporary until automation)	32
4.8 Validation	32
4.8.1 Checking the IPs of the Swarm.....	32
4.8.2 Swarm health and on-instance checks.....	33
4.8.3 Web reachability	33
4.9 CF-Summary.....	33
5. Introducing DynamoDB and reviewing IAM policies.....	33
5.1 Design considerations	34
5.2 We have to begin somewhere. Begin the beginner.....	34
5.2.1 New child template: EC2 role + instance profile	34
5.2.3 Orchestrate this in root.yaml	35
5.2.4 Upload and update the root stack	36
5.2.5 Verify IMDSv2 + role (on the Manager)	37
5.2.6 Summary and an explanation of what we just did.....	38
5.3 Let us introduce DynamoDB	39
5.x Plan to be included:.....	39
6. .NET, .NET, my Java-library for a .NET-app!	39
7. Where to go from here? (Upgrade the system).....	39
X. References.....	40
X.1 educ8.se	40
X.2 LLM support	40

1. Introduction

This is just a tutorial of how we built a webform in .NET that talks through IAM/S3 to AWS DynamoDB and is using a Docker Swarm setup on a scalable AWS EC2-instance setup. I wrote this guide for my own good (learning by describing), but sharing it in case someone else might benefit from it too.

Link to repository: <https://github.com/mymh13/swarm-dotnet-test>

1.1 Pre-requisites

AWS account, AWS CLI installed, SSH key, AWS key pair, Docker (Docker Desktop is handy) and basic cloud and .NET knowledge.

2. Infrastructure basics

In this phase we will do the rudimentary basics: Create a security group for the swarm and launch EC2 instances. Further infrastructure (database) will be added later.

2.1 Create a Security Group

The SG sets the rules and boundaries for the system that uses that particular SG. In this first step, we want to create a SG that primarily handles the inbound rules for our Swarm network.

2.1.1 Create the initial inbound rules

In the AWS console, navigate to EC2 > Security Groups > Create Security Group. Fill in the following:

Name: Choose something descriptive so you know what role the SG plays. You will end up with many SGs eventually and then it will be hard to browse through them.

Description: Something descriptive, see above

VPC: use default VPC for this template, if you can do custom then you do not need this guide 😊

Inbound rules:

Type	Protocol	Port	Source	IP	Description
SSH	TCP	22	My IP	[yourIPhere]	SSH
HTTP	TCP	80	IPv4	0.0.0.0/0	HTTP inbound
Custom TCP	TCP	8080	IPv4	0.0.0.0/0	Visualizer

Create security group Info

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.

Basic details

Security group name Info
Tinfoil-Swarm-SG
Name cannot be edited after creation.

Description Info
Docker Swarm Security Group template

VPC Info
vpc-088ff31b549af5cb0

Inbound rules Info

Type	Protocol	Port range	Source	Description - optional
SSH	TCP	22	My IP	SSH
HTTP	TCP	80	Anywhere...	HTTP inbound
Custom TCP	TCP	8080	Anywhere...	Visualizer

Add rule

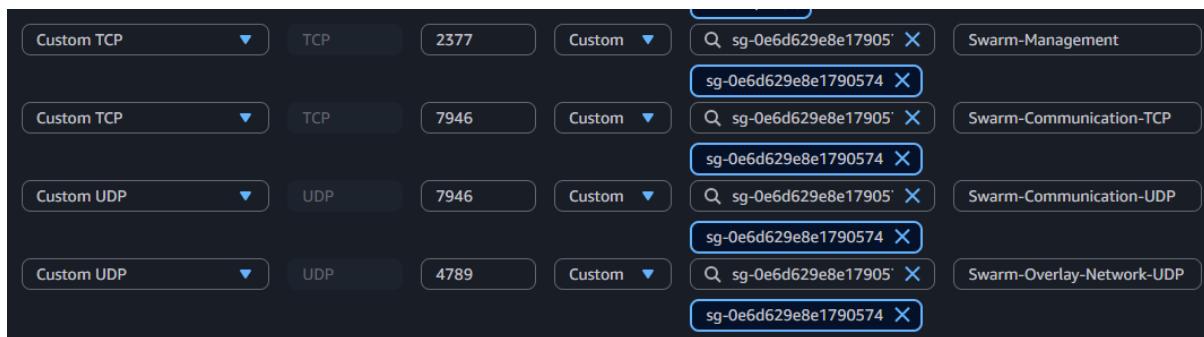
At this point you want to create the security group, there will be four more inbound rules but they will use self-reference so you want to create the SG before it can refer to itself!

2.1.2 Update the Security Group

We need to add the final inbound rules. Edit the Security Group. In the IP/search window, click the window and choice your SG you just created to self-reference.

Inbound rules (3)								
	Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
<input type="checkbox"/>	-	sgr-0e5bb9118013acf06	IPv4	Custom TCP	TCP	8080	0.0.0.0/0	Visualizer
<input type="checkbox"/>	-	sgr-0261101699f05fb43	IPv4	SSH	TCP	22	185.209.198.82/32	SSH
<input type="checkbox"/>	-	sgr-075ac86c81497598d	IPv4	HTTP	TCP	80	0.0.0.0/0	HTTP inbound

Do note UDP on the two last ones! It should look something like this:



2.1.3 What the hell did we just do

Before we move on, let us briefly explain what this does:

- SSH (TCP 22) – Lets you connect to the server via SSH. We lock it down to your own IP only, so no one else can try to log in.
- HTTP (TCP 80) – This allows anyone on the internet to reach our application in the browser.
- Visualizer (TCP 8080) – Docker Swarms visualizer tool runs on port 8080. This rule makes us able to reach and view the visualizer in our browser.
- Swarm Management (TCP 2377) – Used for communication between the swarm manager and the worker nodes. Self-referenced so only other servers in the same SG are allowed in.
- Swarm Communication (TCP 7946 + UDP 7946) – These ports handle internal gossip (yes, that is the actual term in distributed systems! It means each node spread information to its neighbours) and node discovery within the system. Self-reference = self-contained, only swarm members allowed.
- Overlay Network (UDP 4789) – This port carries actual container traffic between swarm nodes over the same overlay network. Self-referenced, internal traffic inside the swarm.

2.1.4 Important note regarding SSH inbound:

Personally, I am always on VPN and my IP changes frequently. This means I have to consider how to handle my SSH IP inbound. In a brief project like this I chose to edit the SSH rule and adjust to reflect my new IP now and then, but for a longer term – and especially if you are more than one person handling this node – you need a more robust solution.

2.1.5 Standard vs custom ports

- Standard: TCP 22 is standard SSH-port. TCP 80 is standard HTTP- (web) traffic.

- Widely used: TCP 8080 is not an official standard, but a widely used convention for dashboards and dev tools.
- Custom: TCP 2377, TCP 7946 + UDP 7946 and UDP 4789 are specific to Docker Swarm.

Custom ports are chosen by the system- or application designer, it is important to know the standards and commonly used ones, so we don't accidentally re-use them for custom roles.

2.1.6 TCP vs UDP protocols

Most internet traffic uses one of the two protocols: TCP or UDP.

- TCP (Transmission Control Protocol) – It sets up a connection, checks that messages arrive, re-sends everything that gets lost. It is reliable but has a bit more overhead. We use this for SSH and HTTP because accuracy is more important than speed.
- UDP (User Datagram Protocol) – It is sent without checking if it arrives, thus making it much faster and uses less overhead than TCP, but delivery is not guaranteed! This is useful for Docker Swarm's internal networking, especially if we are working idempotent as we can accept lost packets.

In short: TCP for reliable traffic, UDP for fast cluster chatter and data transport.

2.1.7 SG-Summary

We have created a Security Group that will handle inbound rules, boundaries and internal traffic.

2.2 Launch EC2 Instances

EC2 Instances is basically just Virtual Machines (VM) but let us stick to AWS-terminology. They will work as "servers" (nodes/hosts might be more appropriate description) for our Swarm. At this point we could also have considered other VMs for roles like a Bastion Host, or a Database or something else, but let us keep this tutorial simple for now and just focus on the Docker Swarm.

2.2.1 First let us create the instances

AWS Console > EC2 > Instances > Launch Instance

Name: Just like with the SG group and classes/methods, always use descriptive names!

AMI: Amazon Linux 2023

Instance type: t3.small

Key pair: This was a pre-req 😊 chose your already pre-set-AWS-key-pair for this

Network settings:

- VPC – default VPC
- Subnet – default Subnet
- Auto-assign public IP – enable
- Firewall / Security Group – Select existing (the one you just created in step 1)

Advanced details > User Data: add a start script for the instance:

```
#!/bin/bash
dnf update -y
dnf install -y docker
```

```
systemctl enable --now docker  
usermod -aG docker ec2-user
```

2.2.1.1 ..and this script does what?

dnf update -y - Updates all existing software on the instance. The -y flag means “say yes to all prompts” so the process doesn’t stop to ask

dnf install -y installs docker - Installs Docker from the Amazon Linux package repository. Again, -y auto-confirms

systemctl enable --now docker - Tells the system to start Docker right away (--now) and to also start it automatically on every reboot (enable)

usermod -aG docker ec2-user - Adds the default AWS login user (ec2-user) to the docker group, so you can run Docker commands without needing sudo every time

2.2.2 Summary, launch and multiples

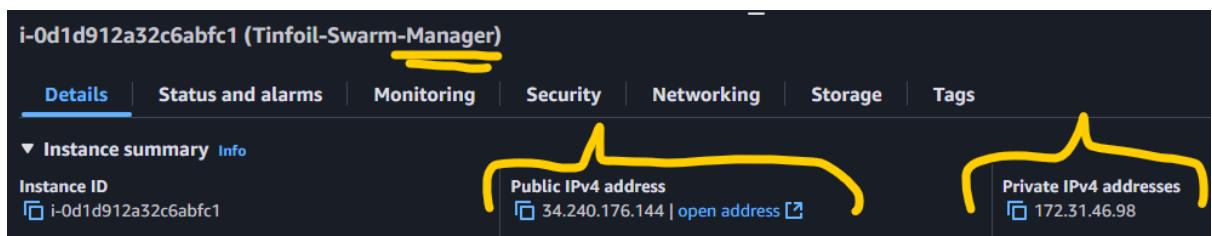
On your right hand, you should have a “Summary” window, with a box named “number of instances” – chose 3 and then click Launch instance.

Then go your EC2 > Instances and rename the second and third Manager-instances to Worker-1 + 2:

Instances (3) Info		Last updated 2 minutes ago	Connect	Instance state ▾	Actions ▾	Launch instances	
Find Instance by attribute or tag (case-sensitive)				All states ▾			
<input type="checkbox"/>	Name 🔗	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
<input type="checkbox"/>	Tinfoil-Swarm-Manager	i-0d1d912a32c6abfc1	Running 🔗 🔍 🔍	t3.small	Initializing 🔗	View alarms +	eu-west-1a
<input type="checkbox"/>	Tinfoil-Swarm-Worker-1	i-03bcacf2d6f7c8f0dc	Running 🔗 🔍 🔍	t3.small	Initializing 🔗	View alarms +	eu-west-1a
<input type="checkbox"/>	Tinfoil-Swarm-Worker-2	i-02c3341722808b0fd	Running 🔗 🔍 🔍	t3.small	Initializing 🔗	View alarms +	eu-west-1a

2.2.3 Note Instance IPs

View the instances (in the version above, September 2025, you can just click the instance tick box to the left and you get a quick view of the instance in a window below the instances). You want to make sure you have the Public IP and the Private IP for the three instances. I copy-pasted them down in a markdown document for easy access.



2.2.4 EC2-Summary

We have set up three EC2 Instances, one working as a Master and two Workers nodes.

2.3 Initialize Docker Swarm via SSH

I hope you have got the IPs handy and SSH + Key Pair is set up. In this phase we will initialize Docker Swarm on our EC2 Instances and verify the cluster. Let's go.

2.3.1 Create a nifty helper-document (optional but recommended)

I used Obsidian and wrote in markdown, but anything will do. What I did was having a number of sections where I listed the public/private IPs of the Instances and a copy-paste command ready to SSH into the Instances. It should look something like this:

```
Swarm Manger
- public IP - 85.482.592.492
- private IP - 172.31.46.98
ssh -i ~/.ssh/<your-key-name>.pem ec2-user@85.482.592.492
```

Do note:

- This assumes your SSH key is in the default user/.ssh/ directory. If you are a Windows user I can highly recommend to always SSH keys in the C:\Users\<username>\.ssh directory
- The public IP in the example above has randomized numbers, this is information you should protect so make sure you do not save this information open in – for example – your project folder (if you use that for a GitHub repo or similar)

2.3.2 Connect to the Manager and initialize Swarm

Use the SSH command [found in 2.3.1](#) to SSH into your Manager node:

```
ssh -i ~/.ssh/your-key.pem ec2-user@<manager-public-ip>
```

Then we will initialize the Swarm on the Manager:

```
sudo docker swarm init --advertise-addr <manager-private-ip>
```

We should get a verification that looks similar to this:

```
Swarm initialized: current node (xyz123) is now a manager.
```

To add a worker to this swarm, run the following command:

```
docker swarm join --token SWMTKN-1-xxx... <manager-private-ip>:2377
```

Copy the join-command and save it to your helper-document.

```
sudo docker swarm join --token SWMTKN-1-<long-string> <manager-private-ip>:2377
```

2.3.2.1 I got most of that. I think

No worries, let me elaborate what you just did.

- By SSHing into the Manager node and running docker swarm init, we turned that server into the “brain” of the swarm.

- The `--advertise-addr <manager-private-ip>` flag makes sure other servers connect over the private AWS network (faster and more secure than public IP).
- Docker then gave us a join token and command. This is like a one-time password that worker nodes will use to join the swarm. Copy it somewhere safe — you'll need it in the next step.

2.3.3 Let us add the Worker nodes

Basically, we will join the existing Swarm as worker nodes, there is not much else to it. Open a new terminal window and run these commands:

```
ssh -i ~/.ssh/your-key.pem ec2-user@<worker1-public-ip>
# Run the join command from step 2.3.2
sudo docker swarm join --token SWMTKN-1-xxx... <manager-private-ip>:2377
```

Repeat the above step but replace `worker1-public-ip` with `worker2`'s public IP. I do not think we need to explain the command in-depth, as you see we use “swarm join” which should be self-explanatory.

You should have three terminal windows that look something like this:

```
~/
[ec2-user@ip-172-31-46-98 ~]$ sudo docker swarm init --advertise-addr 172.31.46.98
Swarm initialized: current node (sfzku1d14tps762bvhufz89t) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-32wel7ks1h0hkee5q3t3virt7d74cjieli9kfivtksfq1nenn-7
    ~~
    ~~.~.
    /_/_/
    _/m/
[ec2-user@ip-172-31-39-19 ~]$ docker swarm join --token SWMTKN-1-32wel7ks1h0hkee5q3t3virt7d74cjieli9kfivtksfq1nenn-7
172.31.46.98:2377
This node joined a swarm as a worker.
[ec2-user@ip-172-31-39-19 ~]$ ...
~~
~~.~.
/_/_/
/_m/
[ec2-user@ip-172-31-42-153 ~]$ docker swarm join --token SWMTKN-1-32wel7ks1h0hkee5q3t3virt7d74cjieli9kfivtksfq1nenn-7
172.31.46.98:2377
This node joined a swarm as a worker.
[ec2-user@ip-172-31-42-153 ~]$ ...
```

2.3.3.1 Let us quickly verify the Swarm cluster

In the Manager nodes' terminal, type this: `sudo docker node ls`

You should see this:

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
7dyzmr cn3ausuqq0mjf8we65f	ip-172-31-39-19.eu-west-1.compute.internal	Ready	Active		25.0.8
1eah1zoejo58fet83gfmk06e	ip-172-31-42-153.eu-west-1.compute.internal	Ready	Active		25.0.8
sfzku1d14tps762bvhufz89t	*	Ready	Active	Leader	25.0.8

2.3.4 Swarm-Summary

We have initialized a Docker Swarm on our Manager, joined with Workers and verified it.

2.4 Deploying the services manually

Ideally this is something that we have an automatic solution for, some IaC solution would be preferable. But this is just a basic template to grasp the concept so we will deploy manually.

2.4.1 Create a Docker Compose file

On the Manager node, we will create the stack file by typing (pasting) this:

```
cat > docker-stack.yml << 'EOF'
version: "3.8"
services:
  web:
    image: nginx:stable-alpine
    deploy:
      replicas: 3
      restart_policy:
        condition: on-failure
      update_config:
        parallelism: 1
        delay: 5s
    ports:
      - "80:80"
    networks: [webnet]

  viz:
    image: dockersamples/visualizer:stable
    deploy:
      placement:
        constraints: [node.role == manager]
    ports:
      - "8080:8080"
    volumes:
      - /var/run/docker.sock:/var/run/docker.sock
    networks: [webnet]

networks:
  webnet:
    driver: overlay
EOF
```

2.4.1.1 Let us not go into all details, but roughly what this does is this

The `cat > docker-stack.yml << 'EOF' ... EOF` command is a way to create a new file (`docker-stack.yml`) and paste content directly into it. Everything between the two EOF markers becomes the file content.

This file is written in YAML, which is very sensitive to spacing and indentation — so it's best to copy-paste rather than type it manually.

The stack defines two services:

- Web: An nginx server with 3 replicas, exposed on port 80
- Viz: A Docker Swarm visualizer, only running on the manager node, exposed on port 8080 (we set that port in the SG but the Manager runs the Viz itself)

At the bottom, it also defines a custom overlay network (webnet) that lets the services talk to each other across swarm nodes.

2.4.2 Deploy the freshly created Stack

Still on the Manager node, run the deploy command:

```
sudo docker stack deploy -c docker-stack.yml myappname
```

```
[ec2-user@ip-172-31-46-98 ~]$ sudo docker stack deploy -c docker-stack.yml myappname
Creating network myappname_webnet
Creating service myappname_viz
Creating service myappname_web
```

This tells Docker Swarm to deploy the stack file (-c docker-stack.yml) and give the stack the name myappname.

2.4.3 Verify the deployment

```
# Show all stacks currently running in the Swarm
sudo docker stack ls

# List the services inside the stack, along with how many replicas are
running. If we followed the guide you should see web with 3/3 replicas and viz
with 1/1 – since we have 3x instances but only the Manager runs the viz.
sudo docker service ls

# Detailed info about each service: which node it is running on, current
state, specific container ID. This is useful when troubleshooting.
sudo docker service ps myappname_web
sudo docker service ps myappname_viz
```

The comment above each command explains them further in-depth.

```
[ec2-user@ip-172-31-46-98 ~]$ sudo docker stack ls
NAME      SERVICES
myappname  2

[ec2-user@ip-172-31-46-98 ~]$ sudo docker service ls
ID        NAME        MODE        REPLICAS   IMAGE                                PORTS
or1o0zokn16  myappname_viz  replicated  1/1      dockersamples/visualizer:stable  *:8080->8080/tcp
oq6uit43aetn  myappname_web  replicated  3/3      nginx:stable-alpine            *:80->80/tcp

[ec2-user@ip-172-31-46-98 ~]$ sudo docker service ps myappname_web
ID        NAME        IMAGE        NODE        DESIRED STATE     CURRENT STATE      ERROR      PORTS
pv1lkvk75478  myappname_web.1  nginx:stable-alpine  ip-172-31-46-98.eu-west-1.compute.internal  Running   Running  3 minutes ago
z78t0hc6xyfj   myappname_web.2  nginx:stable-alpine  ip-172-31-42-153.eu-west-1.compute.internal  Running   Running  3 minutes ago
eg510zadhb03   myappname_web.3  nginx:stable-alpine  ip-172-31-39-19.eu-west-1.compute.internal  Running   Running  3 minutes ago
[ec2-user@ip-172-31-46-98 ~]$ sudo docker service ps myappname_viz
ID        NAME        IMAGE        NODE        DESIRED STATE     CURRENT STATE      ERROR      PORTS
nzat3lxmepib  myappname_viz.1  dockersamples/visualizer:stable  ip-172-31-46-98.eu-west-1.compute.internal  Running   Running  3 minutes ago
```

2.4.4 Deploy-Summary

We have created and deployed the Docker Compose-file on the Manager and verified it.

2.5 Test web service and see if the service can scale

Now let's test that the web service is accessible, and then try scaling it up and down.

2.5.1 Let us first verify web access

Open a browser and go to the public IP of any node:

`http://<public-ip-of-any-node>`

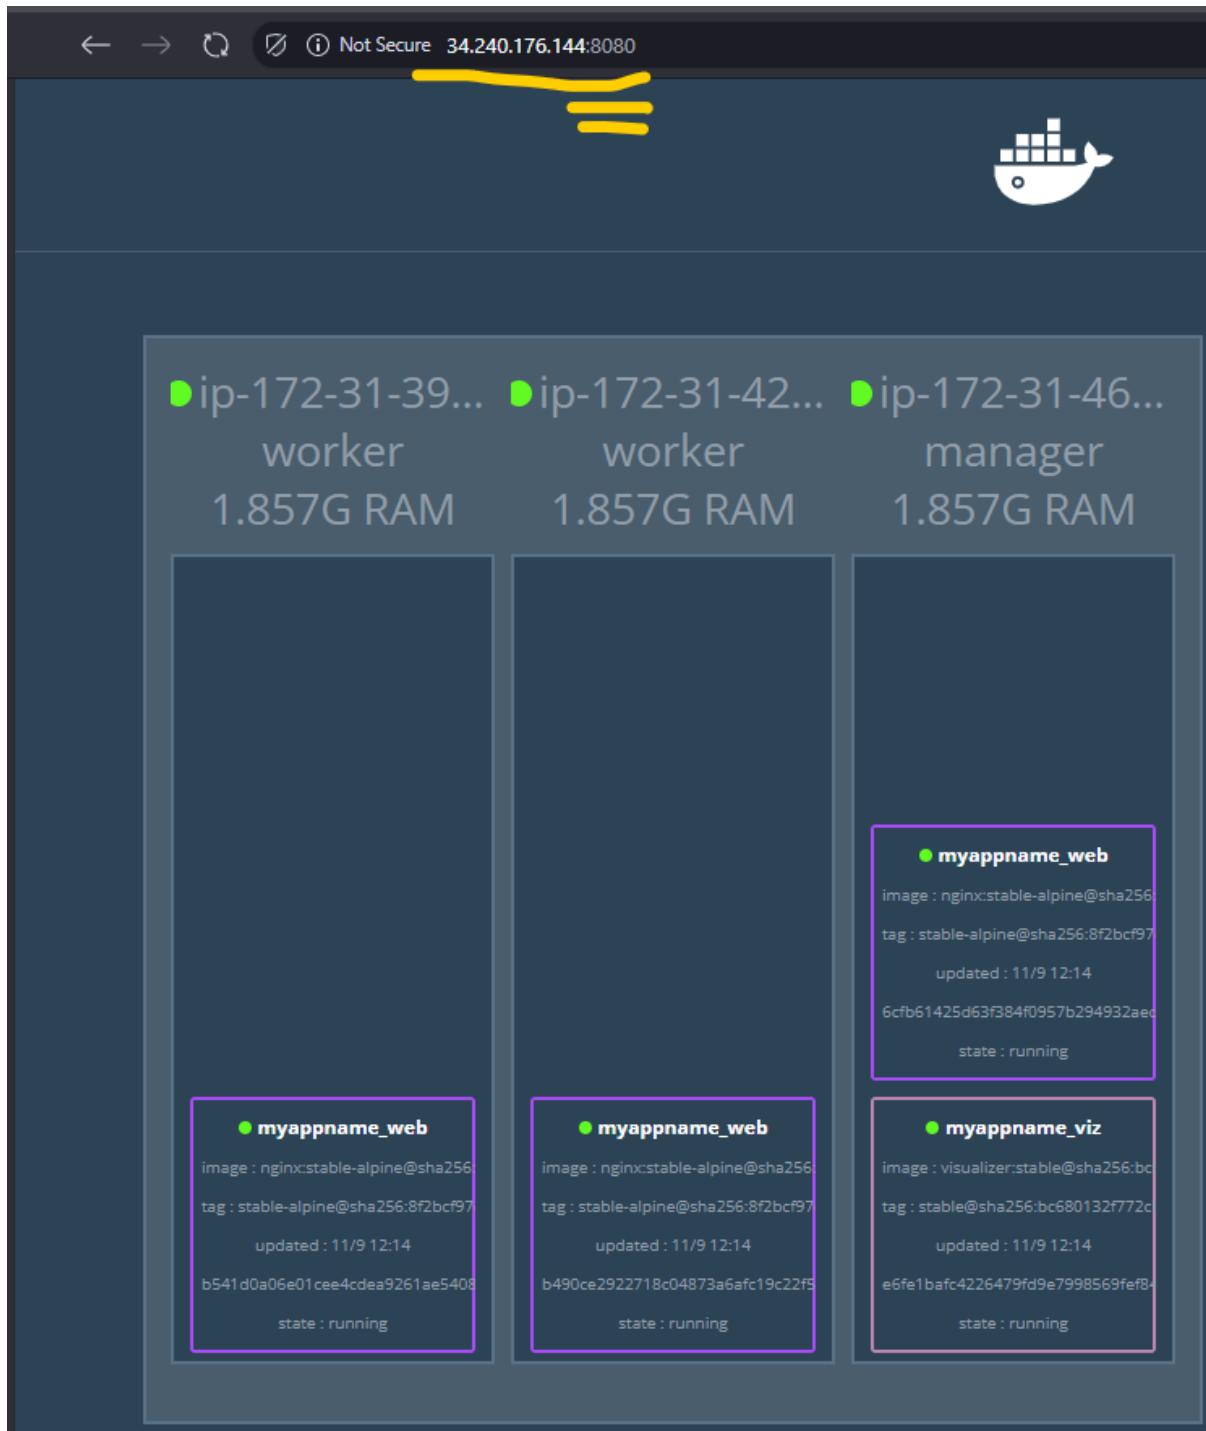
You should see the default Nginx welcome page. Then check the visualizer service:

`http://<manager-public-ip>:8080/`

If everything is configured correctly, you'll see a live visualization of your swarm.

- Port 80 needs to be open for web access.
- Port 8080 needs to be open for the visualizer.

If either page doesn't load, double-check your Security Group and your stack file to make sure the correct ports are open to the right IPs.



2.5.2 Scale the services

Keep that Visualizer browser tab open, and run these commands on the Manager node:

```
# Scale up nginx to 5 replicas
sudo docker service scale myappname_web=5

# Check scaling
sudo docker service ps myappname_web

# Scale back down to 3
```

```
sudo docker service scale myappname_web=3
```

You should see containers appear and disappear in the visualizer as the service scales up and down.

Do note: Scaling here means adding or removing containers, not EC2 instances. Your swarm is still running on the same set of servers. You're just changing how many copies of the nginx container run across them.

2.5.3 Tools to monitor the services

You won't need these right now, as you saw the services running, but leaving this here as it is helpful tools in the future:

```
# Watch service status (Ctrl + C to exit)
watch sudo docker service ls

# View service logs
sudo docker service logs myappname_web
sudo docker service logs myappname_viz
```

“Watch” is live monitoring, so you will have to Ctrl + C to exit / cancel the operation.

Logs are useful tools, as a developer these are your best friends. Logging separate logfiles for separate services / functions is something you will want to do and learn.

2.5.4 Web-Summary

We visit the public IPs, check that the service is running and try the Visualizer.

2.6 Tools for cleaning up (optional)

At this point you might want to clean up a stack (if you are running one), or leave / reset the Swarm.

2.6.1 Remove stack

On the manager node: `sudo docker stack rm myappname`

`rm` means “remove”, it is a good command to know if you are working with files and directories too.

2.6.2 Leave Swarm

You might want to reset the swarm, alter instances or something else:

```
# On workers
sudo docker swarm leave

# On manager (force)
sudo docker swarm leave --force
```

2.6.3 Terminate Instances

Visit EC2 > Instances. Select all three instances. Under “Instance state” chose Terminate.

Instances (3/3) Info

Action

Stop instance
Start instance
Reboot instance
Hibernate instance
Terminate (delete) instance

2.6.4 Delete Security Group

Visit EC2 > Security Groups. Under Actions, Delete the SG.

Actions

View details
Edit inbound rules
Edit outbound rules
Manage tags
Manage stale rules
Copy to new security group
Share security group
Delete security groups

2.7 Summary of above sections

We have built a SG, launched three instances and deployed a Docker Swarm on them, making sure it is healthy, and all the roles are functional.

In the class when we were at this point, we were supposed to do a .NET application and use DynamoDB for handling files that we send (and possibly retrieve, but that was a bonus) from our .NET application. As I have taken down and up these swarms now multiple times, I feel it is appropriate to create section 3 here – create the above system but in CloudFormation (or Terraform, but let us begin with CF).

3 CloudFormation me up, Scotty

I will build this using the Nested Stacks system as I really, strongly dislike messy long files. It has other benefits too:

- Separations of Concern – Each section is independent from other sections. This is good both from a modularity concept but it also easier when trying to solve problems that arise
- Deploy – Deployment benefits greatly from being separated. You can work on one module and just deploy that one, especially if we are working idempotent
- Adaptability – We might want to expand, or reduce, the program. Adapt to various needs. Then it might be beneficial to just replace or remove or add sections that we need.

That said, let us get going!

3.1 S3 Bucket roles first

Before CloudFormation can deploy nested stacks, we need a storage location for the templates and a user that has permissions to interact with that location. This is done in two parts.

3.1.1 Create the bucket first

1. Go to AWS Console > S3 > Create bucket
2. Enter a unique bucket name (ex., cf-artifacts-<account-id>-eu-west-1). Bucket names must be globally unique
3. Select the region where you plan to deploy your stacks (ex., eu-west-1)
4. Leave other settings at defaults for now
5. Click Create bucket

You now have a dedicated bucket for storing CloudFormation templates and artifacts.

3.1.2 Then attach the policy to the user so you can use the bucket

1. Go to AWS Console > IAM > Users > select your user
2. Click Add Permissions > Create inline policy
3. Switch to the JSON tab and paste in your policy (replace <your-bucket-name> with the one you just created in 3.1.1.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3>ListBucket",
        "s3:GetBucketLocation"
      ],
      "Resource": "arn:aws:s3:::<your-bucket-name>"
    },
    {
      "Effect": "Allow",
      "Action": [
        "s3GetObject",
        "s3PutObject",
        "s3DeleteObject",
        "s3DeleteObjectVersion"
      ],
      "Resource": "arn:aws:s3:::<your-bucket-name>/*"
    }
  ]
}
```

Name the policy something descriptive (ex. S3_bucket_policy or cf-s3access-policy) and save it. Your IAM user now has the necessary permissions to interact with the S3 bucket – congrats!

3.2 Quick conversion – CloudFormation style

I recommend visiting the repo for the latest code: <https://github.com/mymh13/swarm-dotnet-test>

Goal: take the manual Swarm setup from sections 1–2 and express it as Infrastructure as Code using nested CloudFormation stacks. We’re not changing what we build—just *how* we build it.

Meow

3.2.1 Nested stacks structure

Child templates live in S3 and a tiny **root** template orchestrates them in order.

```
/infra
  /templates
    root.yaml          # orchestrates children
    00-sg-swarm.yaml   # Security Group for swarm
    10-ec2-swarm.yaml  # EC2 Manager + 2 Workers (Docker installed)
  /parameters
    dev.json           # example parameter set for root.yaml
```

3.2.2 What why when how all those cookies, conceptually? IaC is cool

Mental (music) model:

- root.yaml = the conductor. Look at me waving this tiny stick!
- Child templates = sections of the orchestra (SG, EC2).
- S3 = the sheet music library (where templates live).
- update-stack = the downbeat that makes every-stack play the latest score.

What each piece does (conceptually):

- root.yaml (master): Orchestrates order and wiring. It calls the SG stack first, then the EC2 stack, and passes outputs (like the SG ID) to the next stack’s inputs. It’s the single command you use to create, update, and delete *everything*.
- 00-sg-swarm.yaml (security): Encodes the exact ports you opened manually (SSH 22 locked to your IP, HTTP 80, Visualizer 8080 to world; Swarm mgmt/overlay ports self-referenced). Same rules, now declarative.
- 10-ec2-swarm.yaml (compute): Launches three EC2s (1 manager, 2 workers) in your chosen subnet/SG and runs user-data to install Docker. This mirrors your earlier “install Docker on each node” step.

Why nested stacks for this:

- Separation of concerns: Each child template does one thing well. Easier to reason about, change, and reuse.
- Safer iteration: You can update just one child (ex., SG rules) and then update-stack the root—CloudFormation figures out the minimal change.
- Single tear-down: Deleting the root stack removes the whole setup in the right order.
- S3 as source of truth: You upload the child templates to S3; root.yaml references them via TemplateURL. When you push new child templates to S3, the next update-stack picks them up.

What we do not automate (yet):

We stop at “Docker installed.” You still:

- SSH to the manager > docker swarm init --advertise-addr <manager-private-ip>
- SSH to each worker > run the docker swarm join ... the manager printed
- Deploy docker-stack.yml as before
(We'll automate init/join later with roles/SSM or cloud-init once the foundation is stable.)

Parameters & environments (lightweight, IaC-friendly):

We keep a small infra/parameters/dev.json (local, ignored by Git) with things like:

- TemplateBucket, TemplatePrefix (where child templates live in S3)
- VpcId, PublicSubnetId (use default VPC/subnet for now)
- AllowedSshCidr, KeyPairName, InstanceType, Amild

This makes create/update/teardown a single command each (no ad-hoc clicking, no scripts).

3.2.3 root.yaml (Master stack)

This will be a long “skeleton” code, and this all assumes you have basic understanding of AWS and Cloudformation. The SG/EC2-setup we run should be self-explanatory more or less though. Template:

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Root stack for Docker Swarm (manager + 2 workers) using nested stacks

Parameters:
  StackNamePrefix:
    Type: String
    Default: swarm-cf
    Description: Prefix used for naming
  TemplateBucket:
    Type: String
    Description: S3 bucket that stores child templates (e.g. cf-artifacts-<acct>-<region>)
  TemplatePrefix:
    Type: String
    Default: swarm-iac/templates/
    Description: Key prefix inside the S3 bucket
  VpcId:
    Type: AWS::EC2::VPC::Id
    Description: VPC to deploy into (choose your default VPC)
  PublicSubnetId:
    Type: AWS::EC2::Subnet::Id
    Description: Public subnet for all three instances (default subnet is fine)
  AllowedSshCidr:
```

```

Type: String
Description: Your IP in CIDR for SSH (e.g. 1.2.3.4/32)
KeyPairName:
  Type: AWS::EC2::KeyPair::KeyName
  Description: Existing EC2 key pair name
InstanceType:
  Type: String
  Default: t3.small
AmiId:
  Type: AWS::EC2::Image::Id
  Description: Amazon Linux 2023 AMI in this region

Resources:
SwarmSecurityGroupStack:
  Type: AWS::CloudFormation::Stack
  Properties:
    TemplateURL: !Sub
      'https://s3.${AWS::Region}.amazonaws.com/${TemplateBucket}/${TemplatePrefix}00
      -sg-swarm.yaml'
    Parameters:
      StackNamePrefix: !Ref StackNamePrefix
      VpcId: !Ref VpcId
      AllowedSshCidr: !Ref AllowedSshCidr

SwarmEc2Stack:
  Type: AWS::CloudFormation::Stack
  DependsOn: SwarmSecurityGroupStack
  Properties:
    TemplateURL: !Sub
      'https://s3.${AWS::Region}.amazonaws.com/${TemplateBucket}/${TemplatePrefix}10
      -ec2-swarm.yaml'
    Parameters:
      StackNamePrefix: !Ref StackNamePrefix
      PublicSubnetId: !Ref PublicSubnetId
      SecurityGroupId: !GetAtt
        SwarmSecurityGroupStack.Outputs.SwarmSecurityGroupId
      KeyPairName: !Ref KeyPairName
      InstanceType: !Ref InstanceType
      AmiId: !Ref AmiId

Outputs:
  ManagerPublicIp:
    Value: !GetAtt SwarmEc2Stack.Outputs.ManagerPublicIp
  Worker1PublicIp:
    Value: !GetAtt SwarmEc2Stack.Outputs.Worker1PublicIp
  Worker2PublicIp:
    Value: !GetAtt SwarmEc2Stack.Outputs.Worker2PublicIp
  SecurityGroupId:

```

```
Value: !GetAtt SwarmSecurityGroupStack.Outputs.SwarmSecurityGroupId
```

3.2.4 00-swarm.yaml (Swarm Security Group)

This stack controls the Swarms SG.

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Security Group for Swarm (SSH, HTTP, Viz, Swarm ports)

Parameters:
  StackNamePrefix:
    Type: String
  VpcId:
    Type: AWS::EC2::VPC::Id
  AllowedSshCidr:
    Type: String

Resources:
  SwarmSG:
    Type: AWS::EC2::SecurityGroup
    Properties:
      GroupDescription: !Sub '${StackNamePrefix}-swarm-sg'
      VpcId: !Ref VpcId
      # Only non-self-referencing rules here
      SecurityGroupIngress:
        # SSH (locked to your IP)
        - IpProtocol: tcp
          FromPort: 22
          ToPort: 22
          CidrIp: !Ref AllowedSshCidr
        # HTTP 80 (world)
        - IpProtocol: tcp
          FromPort: 80
          ToPort: 80
          CidrIp: 0.0.0.0/0
        # Visualizer 8080 (world)
        - IpProtocol: tcp
          FromPort: 8080
          ToPort: 8080
          CidrIp: 0.0.0.0/0
      # Egress is outward traffic
      SecurityGroupEgress:
        - IpProtocol: -1
          CidrIp: 0.0.0.0/0
    Tags:
      - Key: Name
        Value: !Sub '${StackNamePrefix}-swarm-sg'
```

```

# Self-referencing rules as separate resources (avoids circular dependency)
IngressSwarmMgmt2377:
  Type: AWS::EC2::SecurityGroupIngress
  Properties:
    GroupId: !Ref SwarmSG
    IpProtocol: tcp
    FromPort: 2377
    ToPort: 2377
    SourceSecurityGroupId: !Ref SwarmSG

IngressGossipTCP7946:
  Type: AWS::EC2::SecurityGroupIngress
  Properties:
    GroupId: !Ref SwarmSG
    IpProtocol: tcp
    FromPort: 7946
    ToPort: 7946
    SourceSecurityGroupId: !Ref SwarmSG

IngressGossipUDP7946:
  Type: AWS::EC2::SecurityGroupIngress
  Properties:
    GroupId: !Ref SwarmSG
    IpProtocol: udp
    FromPort: 7946
    ToPort: 7946
    SourceSecurityGroupId: !Ref SwarmSG

IngressOverlayUDP4789:
  Type: AWS::EC2::SecurityGroupIngress
  Properties:
    GroupId: !Ref SwarmSG
    IpProtocol: udp
    FromPort: 4789
    ToPort: 4789
    SourceSecurityGroupId: !Ref SwarmSG

Outputs:
  SwarmSecurityGroupId:
    Value: !Ref SwarmSG

```

3.2.5 10-ec2-swarm.yaml (EC2 Manager + 2 Workers, Docker installed)

Here we create the EC2-Swarm, and set the User Data to install docker, like we did [here](#).

```

AWSTemplateFormatVersion: '2010-09-09'
Description: EC2 instances for Swarm (1 manager, 2 workers) with Docker
installed

```

```

Parameters:
  StackNamePrefix:
    Type: String
  PublicSubnetId:
    Type: AWS::EC2::Subnet::Id
  SecurityGroupId:
    Type: String
  KeyPairName:
    Type: AWS::EC2::KeyPair::KeyName
  InstanceType:
    Type: String
  AmiId:
    Type: AWS::EC2::Image::Id

Mappings: {}

Resources:
  Manager:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: !Ref AmiId
      InstanceType: !Ref InstanceType
      KeyName: !Ref KeyPairName
      SubnetId: !Ref PublicSubnetId
      SecurityGroupIds: [ !Ref SecurityGroupId ]
      Tags:
        - Key: Name
          Value: !Sub '${StackNamePrefix}-manager'
    UserData:
      Fn::Base64: !Sub |
        #cloud-config
        runcmd:
          - dnf update -y
          - dnf install -y docker
          - systemctl enable --now docker
          - usermod -aG docker ec2-user

  Worker1:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: !Ref AmiId
      InstanceType: !Ref InstanceType
      KeyName: !Ref KeyPairName
      SubnetId: !Ref PublicSubnetId
      SecurityGroupIds: [ !Ref SecurityGroupId ]
      Tags:
        - Key: Name

```

```

    Value: !Sub '${StackNamePrefix}-worker-1'
  UserData:
    Fn::Base64: !Sub |
      #cloud-config
      runcmd:
        - dnf update -y
        - dnf install -y docker
        - systemctl enable --now docker
        - usermod -aG docker ec2-user

Worker2:
  Type: AWS::EC2::Instance
  Properties:
    ImageId: !Ref AmiId
    InstanceType: !Ref InstanceType
    KeyName: !Ref KeyPairName
    SubnetId: !Ref PublicSubnetId
    SecurityGroupIds: [ !Ref SecurityGroupId ]
  Tags:
    - Key: Name
      Value: !Sub '${StackNamePrefix}-worker-2'
  UserData:
    Fn::Base64: !Sub |
      #cloud-config
      runcmd:
        - dnf update -y
        - dnf install -y docker
        - systemctl enable --now docker
        - usermod -aG docker ec2-user

Outputs:
  ManagerPublicIp:
    Value: !GetAtt Manager.PublicIp
  Worker1PublicIp:
    Value: !GetAtt Worker1.PublicIp
  Worker2PublicIp:
    Value: !GetAtt Worker2.PublicIp

```

Note:

This intentionally stops at Docker installed. After stack completes, follow the tutorial's SSH steps:

- SSH to manager, run docker swarm init --advertise-addr <manager-private-ip>.
- SSH to each worker, run the docker swarm join ... command from the manager output.
- Deploy your docker-stack.yml exactly as before.

3.2.6 Nested stacks – big picture (TL;DR)

- You upload child templates to S3.
- root.yaml points to those S3 URLs and wires stacks in order.

- Create/Update the *root* stack only > everything else follows.
- Delete the *root* stack > everything tears down cleanly.
- Manual Swarm init/join stays for now (keeps the foundation simple); we'll automate later.

3.3 Important note on SSH, dynamic IP ranges, ingress/egress

If you like me run a VPN that alternates your IP now and then, it can be problematic to run “your own IP” as SSH-in-setting as per the [2.1.1 Inbound rules](#) we set.

3.3.1 This is important to know re: SSH / AllowedSshCidr

- The template does not auto-detect your IP. AllowedSshCidr is a parameter you set in dev.params.json (or via CLI) before create-stack / update-stack.
- If your VPN gives you a new IP, you'll need to edit AllowedSshCidr to <new-ip>/32 and run update-stack to open SSH again.
- Existing SSH sessions remain allowed as long as your source IP doesn't change. If the VPN reassigns your IP mid-session, the TCP connection will drop and you'll need to update the rule and reconnect.
- Alternatives (for later):
 - Use a stable CIDR from your VPN provider (if they publish one) instead of a single /32
 - SSM Session Manager (no inbound 22 at all) by attaching an IAM role and enabling SSM on the instances
 - A bastion host with a fixed IP/Security Group that's allowed to SSH, while nodes deny world-wide SSH

So your dev.params.json probably have this, and you need to modify it each time you swap IP:

```
{
  "ParameterKey": "AllowedSshCidr",
  "ParameterValue": "your.id.range.here/32"
},
```

You can find your IP by running this:

```
curl -s https://checkip.amazonaws.com
```

3.3.2 Let us also mention egress rules

- SecurityGroupEgress: - IpProtocol: -1, CidrIp: 0.0.0.0/0 = allow all outbound.
- Security groups are stateful, so responses to allowed inbound are automatically allowed out; however:
 - Your instances need to initiate outbound connections (e.g., dnf update, pulling Docker images) – important!
 - That requires an egress allow. The default is “allow all egress”; we set it explicitly for clarity. Technically it is not needed but this shows intent
- Keep it as-is for this public-subnet setup. (When/if you move to private subnets + NAT, the rule still makes sense—traffic exits via NAT. I.e good habit to declare it)

3.4 CF-Summary

Hopefully we have successfully converted our manual deploy to CloudFormation templates now. I explicitly tried not to paste a lot of code in here so you (the reader) would just copy-paste, I try to explain what the code does and what. That is a bit hard to do with the CloudFormation templates, so I just assume you know the CF and AWS basics.

A bit short explanation would probably be that CF declare parameters (like the ingredients if we are reading a cooking recipe) and then we specify the instructions how to cook (resources, etc). Finally we share that information through the Outputs section.

4. Deploying our CF and validating it

We keep this fully IaC + CLI-driven: validate locally, upload the child templates to S3, then create/update/delete the single root stack.

4.1 Tiny quality of life life-hacks

We will do two sections below, the second is mandatory, the first is optional: but it makes life so much better that it will be worth setting that too.

4.1.1 Optional, tiny .env for copy-pasting shorter CLI

This .env ends up in the /infra/ directory, here it will be ignored by .gitignore. Create that .env file and add this code to it:

```
# .env (example)
export REGION="eu-west-1"
export BUCKET="cf-swarm-<user-id-here>-eu-west-1"
export PREFIX="swarm-iac/templates/"
export STACK="swarm-cf-root"
export PARAMS="infra/parameters/dev.params.json"
```

To use (export) those values in your shell session, type the following in Bash (and make sure you are on the same level as the .env file, I keep mine in the Git-root as that is usually my default):

```
source .env
```

If you want to verify that this worked (not a bad idea), run this:

```
for v in REGION BUCKET PREFIX STACK PARAMS; do echo "$v=${!v}"; done
```

It will print out the stored parameter values in your console so you can see if they are correct.

4.1.2 Create dev.params.json

Our json template needs to be in an array form for CloudFormation to be able to read it. I handled this by creating two new files: dev.params.json and dev.example.json. The example-version show

what it might look like. While the params-version is the real deal, gitignored. While at it, delete the dev.json template we had as it was replaced by the above versions.

It is time to set the real values to dev.params.json so we can use those in the system. For obvious reasons I cannot post my personal information here, just replace the data with your own:

- StackNamePrefix — it is there to add a prefix to names we create from templates
- TemplateBucket/TemplatePrefix — where child templates live in S3: path routes
- VpcId/PublicSubnetId — we're using your existing (default) VPC; just point to one public subnet
- AllowedSshCidr — your IP in /32 CIDR for SSH (port 22) to the instances
- KeyPairName — existing EC2 key pair (no .pem)
- InstanceType — size for the EC2's
- AmiId — the AMI to boot (Amazon Linux 2023 in this case), you can find this when you browse the AMI alternatives (but for simplicity: use ami-097f734cebd08c39e)

This below is the example file, some data is correct but replace the rest:

```
[  
  { "ParameterKey": "StackNamePrefix", "ParameterValue": "swarm-cf" },  
  { "ParameterKey": "TemplateBucket", "ParameterValue": "cf-artifacts-<account-id>-eu-north-1" },  
  { "ParameterKey": "TemplatePrefix", "ParameterValue": "swarm-iac/templates/" },  
  { "ParameterKey": "VpcId", "ParameterValue": "vpc-xxxxxx" },  
  { "ParameterKey": "PublicSubnetId", "ParameterValue": "subnet-xxxxxx" },  
  { "ParameterKey": "AllowedSshCidr", "ParameterValue": "1.2.3.4/32" },  
  { "ParameterKey": "KeyPairName", "ParameterValue": "your-keypair" },  
  { "ParameterKey": "InstanceType", "ParameterValue": "t3.small" },  
  { "ParameterKey": "AmiId", "ParameterValue": "ami-097f734cebd08c39e" }  
]
```

4.1.2.1 *Important note, do not forget to set values*

At this point, I did the rookie mistake by forgetting to set the values for the templates I gave you above. I forgot to replace the xxxx after vpc and subnet with their actual values.

4.1.2.2 *Where to find the VPC ID*

VPC can be found at AWS Console > VPC > Your VPCs – just copy the ID and replace in the file.

4.1.2.3 *Where to find the Subnet ID and which to pick and why*

Subnets can be found in the menu just under “Your VPCs”, the step we just visited. You will see three default Subnets, so to know which to pick you need to understand Subnet basics. You see this:

- Default VPC CIDR: 172.31.0.0/16 (a /16 = 65,536 IPs).
- Default subnets (one per AZ) like:
 - 172.31.0.0/20
 - 172.31.16.0/20
 - 172.31.32.0/20

A /20 is 4,096 IPs. It's a slice of the /16. The “16” and “32” offsets are just the next /20 blocks (each /20 advances by 16 in the third octet for 172.31.x.0). In the default VPC, these default subnets are public (they auto-assign public IPs and have a route to the Internet Gateway).

What to pick now?

- Pick any one default subnet in that VPC where MapPublicIpOnLaunch = true. All three are usually public in the default VPC—choose one.
- Why one subnet is fine (for now)
- Our minimal template places all 3 instances in one public subnet for simplicity. We'll spread across AZs later when we add ALB/ASG, etc.

A summary would be that the Subnets will be useful when you need to extend IP ranges and create internal solutions. We have a rudimentary network setup so we don't need that type of internal communication, we are just solving it by opening ports and using self-reference for the Swarm.

4.1.2.4 Using CLI to solve the above

I am not going in-depths on these commands, what they do is to search for the information you were looking for manually in the steps above:

```
# 1) Default VPC in your region
aws ec2 describe-vpcs \
--region "$REGION" \
--filters Name=isDefault,Values=true \
--query 'Vpcs[0].VpcId' --output text

# 2) List subnets in that VPC; pick any with PublicIpOnLaunch = true
VPC_ID="vpc-xxxxxxxx" # <- paste from the previous command
aws ec2 describe-subnets \
--region "$REGION" \
--filters Name=vpn-id,Values=$VPC_ID \
--query
'Subnets[].{SubnetId:SubnetId,Az:AvailabilityZone,Cidr:CidrBlock,PublicIpOnLaunch:MapPublicIpOnLaunch}' \
--output table
```

4.2 Validate template (syntax check)

- You've created the artifacts S3 bucket and given your user S3 permissions ([section 3.1](#))
- AWS CLI is configured for the target account
- Your three templates exist:
 1. infra/templates/00-sg-swarm.yaml
 2. infra/templates/10-ec2-swarm.yaml
 3. infra/templates/root.yaml

So now we will bring forward one of our strongest weapons: a simple syntax validation check. We will run this after we make updates to our CloudFormation YAMLS, it will save us time and headache:

```
aws cloudformation validate-template --template-body file://infra/templates/00-sg-swarm.yaml
aws cloudformation validate-template --template-body file://infra/templates/10-ec2-swarm.yaml
aws cloudformation validate-template --template-body file://infra/templates/root.yaml
```

Pardon the small font, I did not want them to be cut off, it is worth listing them on a line like that. My point is this: you will want to copy paste these babies frequently. I will link to this section.

Be mindful that these validate-template commands check syntax, it does not fetch child TemplateURLs. This is why we have to do individual validations, otherwise we could have used root.

4.3 Upload the templates to S3

This is another command you will run frequently. Upload the files to S3 – child firsts, then root.

```
aws s3 cp infra/templates/00-sg-swarm.yaml "s3://${BUCKET}/${PREFIX}00-sg-swarm.yaml"
aws s3 cp infra/templates/10-ec2-swarm.yaml "s3://${BUCKET}/${PREFIX}10-ec2-swarm.yaml"
aws s3 cp infra/templates/root.yaml "s3://${BUCKET}/${PREFIX}root.yaml"
```

What this does: cp copies the file from path/filename.yaml and then it ships it to s3://the-bucket-variable-we-defined/(adds-the-prefix-to-our-filename)filename.yaml. root will have the templateURL inside S3 for the child templates, this is why it is important to upload the childs first so you have the latest version pulled when you run update-stack on root.

4.4 Create the stack (daily bring-up of the stacks)

This is a command you will run everytime you want to bring up the (most recently uploaded) stacks on AWS. For these school/non-live versions I would suggest you tear down the project daily and then you re-deploy it with this command on a daily basis. Since a new session most likely will occur in your shell (Bash) when you start up, remember to run the export command first from 4.1.1. Otherwise you will need to adjust these commands too since we parametrized values! Then run this:

```
aws cloudformation create-stack \
--region "$REGION" \
--stack-name "$STACK" \
--template-url
"https://s3.${REGION}.amazonaws.com/${BUCKET}/${PREFIX}root.yaml" \
--parameters file://"${PARAMS}" \
--capabilities CAPABILITY_IAM CAPABILITY_NAMED_IAM

aws cloudformation wait stack-create-complete \
--region "$REGION" --stack-name "$STACK"
```

Create stack, in the region and stack-name you predetermined. Template-url will be the path given, parameters are taken from \$PARAMS, capabilities IAM. Maybe I used this command too many times, if you feel it needs further explanation please reach out to me, it seems rather straightforward to me.

Verification: you should get a return prompt that gives you the Stack ID.

4.4.1 Grab outputs (IPs, SG id)

Run this command to get a table printout of the values – since you do get a new IP and SG-ID daily:

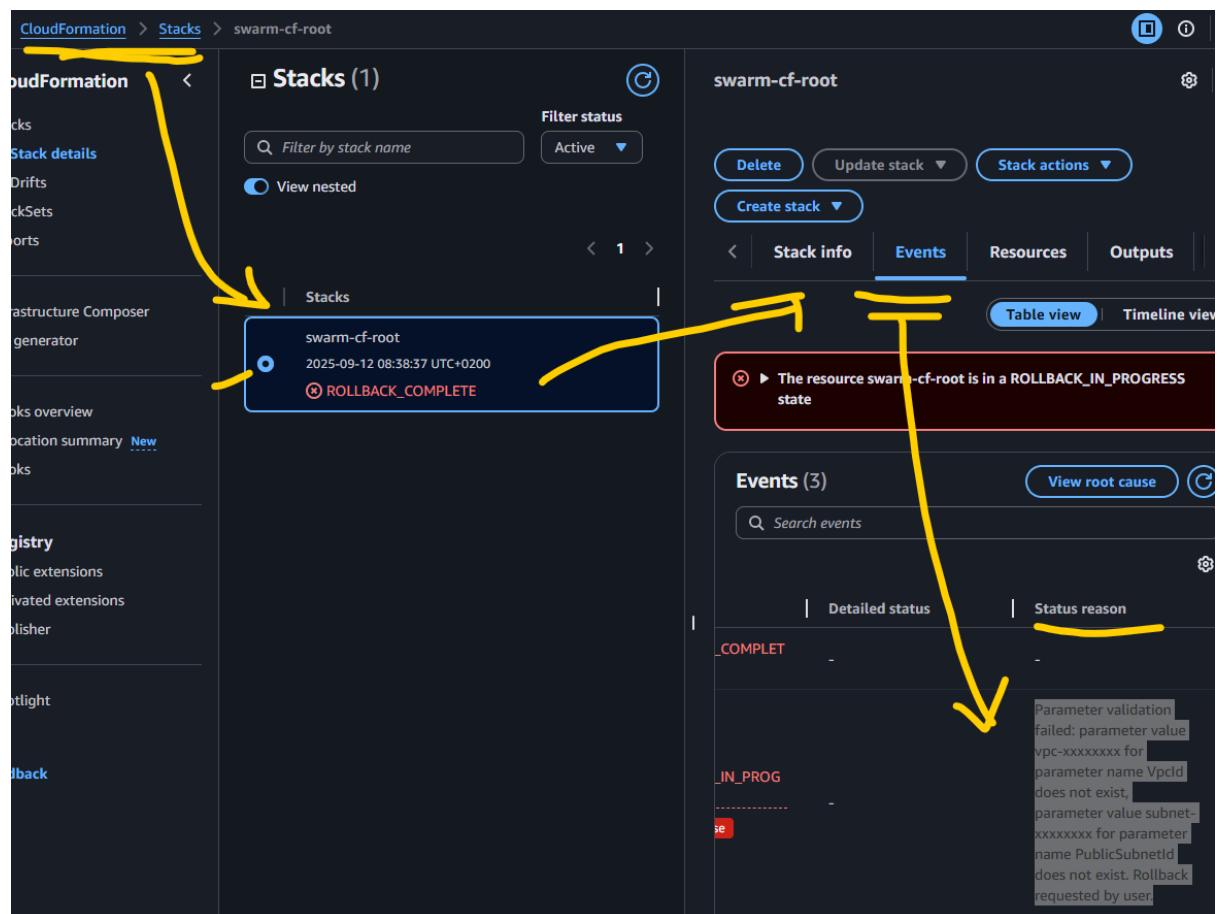
```
aws cloudformation describe-stacks \
--region "$REGION" --stack-name "$STACK" \
--query 'Stacks[0].Outputs[].[OutputKey,OutputValue]' --output table
```

I go more in-depth on this topic at [section 4.7](#).

4.4.2 Error handling when creating

Sometimes you do something naughty and you will get slapped on the fingers. The good news is the system is good at telling you what you did wrong. Logs are your best friends. One such example could be that you forgot to replace the “vpc-xxxxxxxxxx” parameter in dev.params.json, like I did. But you can find the error easily.

When you upload your files (and this is why I put this error handling under section 4.4), you will get a “ROLLBACK_COMPLETE” or similar error in your CloudFormation console at AWS. If you click your stack there, and check the Events tab, it will provide information about what you did wrong. Like this:



In this case the solution was easy: just apply the proper information in dev.params.json, re-upload the root.yaml and re-run stack creation/update (depending where you are in your process).

4.5 Update the stack (apply template changes)

Whenever you have updated templates, [validate](#), [re-upload it to S3](#), then run update-stack like this:

```

aws cloudformation update-stack \
--region "$REGION" \
--stack-name "$STACK" \
--template-url
"https://s3.${REGION}.amazonaws.com/${BUCKET}/${PREFIX}root.yaml" \
--parameters file://${PARAMS} \
--capabilities CAPABILITY_IAM CAPABILITY_NAMED_IAM

aws cloudformation wait stack-update-complete \
--region "$REGION" --stack-name "$STACK"

```

It is a similar version to the create-stack, but instead it updates the stack. Remember: root last!

4.6 Tear-down of the stacks (daily or frequent routine)

When you are done working and take a break, tear down the stack(s) (you tear it down using the STACK variable which is a mutual name for the primary stack that root commands, thus all child templates go down with it) by using this command:

```

aws cloudformation delete-stack \
--region "$REGION" \
--stack-name "$STACK"

aws cloudformation wait stack-delete-complete \
--region "$REGION" --stack-name "$STACK"

```

4.7 Optional: Check the public IPs of the Instances

We will need the public IP for the temporary step in 4.7.1, but it also nice to know it if we would need to SSH in or whatever. Here is how you do that:

```

aws cloudformation describe-stacks \
--region "$REGION" --stack-name "$STACK" \
--query 'Stacks[0].Outputs[].[OutputKey,OutputValue]' --output table

```

```

$ aws cloudformation describe-stacks \
--region "$REGION" --stack-name "$STACK" \
--query 'Stacks[0].Outputs[].[OutputKey,OutputValue]' --output table
-----
|             DescribeStacks           |
+-----+
| Worker1PublicIp | 108.130.168.64   |
| Worker2PublicIp | 63.34.171.184    |
| SecurityGroupId | sg-02167463e11a80d6c |
| ManagerPublicIp | 54.78.33.150    |
+-----+

```

Quite nice! And useful. 😊

4.7.1 Bonus: Capture the three public IPs into variables

Oh snap now we are into bonus-land-deluxe, but developers are supposed to automate..

```
MANAGER_PUB=$(aws cloudformation describe-stacks --region "$REGION" --stack-name "$STACK" \
    --query "Stacks[0].Outputs[?OutputKey=='ManagerPublicIp'].OutputValue" --output text)
WORKER1_PUB=$(aws cloudformation describe-stacks --region "$REGION" --stack-name "$STACK" \
    --query "Stacks[0].Outputs[?OutputKey=='Worker1PublicIp'].OutputValue" --output text)
WORKER2_PUB=$(aws cloudformation describe-stacks --region "$REGION" --stack-name "$STACK" \
    --query "Stacks[0].Outputs[?OutputKey=='Worker2PublicIp'].OutputValue" --output text)

echo "Manager: $MANAGER_PUB"
echo "Worker1: $WORKER1_PUB"
echo "Worker2: $WORKER2_PUB"
```

This grabs the three public IPs and stores them into the variables. See 4.7.2 for :usefulness:

4.7.2 Temporary: Until we automate deploy, do it manually

1. SSH to Manager > init swarm. See [section 2.3.2](#) and follow those steps.
2. SSH to each worker, same link as above.
3. Deploy the stack from the Manager, same as above too.

If you were a geek and did step 4.7.1 for the bonus-capture you can do like this:

```
# Manager
ssh -i ~/.ssh/tinfoil-eu-west-1.pem ec2-user@$MANAGER_PUB
sudo docker swarm init --advertise-addr "$MANAGER_PRIV"
# Copy the printed join command
# Type "exit" to leave this node

# Worker 1
ssh -i ~/.ssh/tinfoil-eu-west-1.pem ec2-user@$WORKER1_PUB
sudo docker swarm join --token SWMTKN-1-... "$MANAGER_PRIV:2377"
# Obviously, replace the ... with the actual join command
# exit

# Worker 2
ssh -i ~/.ssh/tinfoil-eu-west-1.pem ec2-user@$WORKER2_PUB
sudo docker swarm join --token SWMTKN-1-... "$MANAGER_PRIV:2377"
```

That is pretty cool. And lazy. Which we like. No more copy/pasting or writing down IPs.

4.7.3 Deploying the Docker Compose (temporary until automation)

We keep this step manual as in [step 2.4.1](#) and [2.4.2](#). SSH into the Manager, create the compose file (as in 2.4.1), deploy it as in 2.4.2.

4.8 Validation

There are all kinds of ways to validate, I believe in validating small incremental steps as you build, if you validate too much at this step then there is a decent risk that you have a lot of issues created along the way that is a lot more messy to solve now. But, for the sake of validation, I will show you a number of tools we have at our disposal here:

4.8.1 Checking the IPs of the Swarm

Run this command in Bash:

```
MANAGER_PUB=$(aws cloudformation describe-stacks --region "$REGION" --stack-name "$STACK" \
    --query "Stacks[0].Outputs[?OutputKey=='ManagerPublicIp'].OutputValue" \
    --output text)
WORKER1_PUB=$(aws cloudformation describe-stacks --region "$REGION" --stack-name "$STACK" \
    --query "Stacks[0].Outputs[?OutputKey=='Worker1PublicIp'].OutputValue" \
    --output text)
WORKER2_PUB=$(aws cloudformation describe-stacks --region "$REGION" --stack-name "$STACK" \
    --query "Stacks[0].Outputs[?OutputKey=='Worker2PublicIp'].OutputValue" \
    --output text)

aws ec2 describe-instances --region "$REGION" \
    --filters "Name=ip-
address,Values=${MANAGER_PUB},${WORKER1_PUB},${WORKER2_PUB}" \
    --query 'Reservations[].Instances[][Tags[?Key==`Name`][0].Value,PublicIpAddress,PrivateIpAddress]' --output table
```

You will get an output that looks like this:

```
aws ec2 describe-instances --region "$REGION" \
    --filters "Name=ip-address,Values=${MANAGER_PUB},${WORKER1_PUB},${WORKER2_PUB}" \
    --query 'Reservations[].Instances[][Tags[?Key==`Name`][0].Value,PublicIpAddress,PrivateIpAddress]' --output table
-----
|          DescribeInstances          |
+-----+-----+-----+
| swarm-cf-worker-1 | 108.130.168.64 | 172.31.10.241 |
| swarm-cf-manager | 54.78.33.150  | 172.31.2.233  |
| swarm-cf-worker-2 | 63.34.171.184 | 172.31.13.18  |
+-----+-----+-----+
```

To the left is your Swarm nodes and their names. In the middle you have the public IP, the one you can visit to check these instances in a browser. To the right is their private IP. Notice how they correlate with the Subnet you set earlier.

4.8.2 Swarm health and on-instance checks

Like we did in [2.4.3](#). Since you are doing those commands on the manager, you can also do on-instance checks from the manager. This is overkill. But hey. Better be safe than sorry? 😊

```
sudo ss -tulpn | egrep ':80|:8080'    # ports open
sudo docker logs $(sudo docker ps -q --filter name=myappname_web -n 1) --
tail=50
```

4.8.3 Web reachability

I think you should do this two ways. First through shell / Bash:

```
# Nginx on any node (HTTP 200 OK)
curl -I http://$MANAGER_PUB
curl -I http://$WORKER1_PUB
curl -I http://$WORKER2_PUB
```

This gives you more detailed info about the state of the machines, but what you want to see primarily is a HTTP/1.1 200 OK at the start. It is nice to see the Nginx server and the Content-type: text-html there too, this means your html page is displaying properly.

Then as step two just visit the public IPs in a browser. Don't forget to test the managers public IP with port :8080 and see if you can see the Visualizer, it should display, then you know the SGs port setting for the Visualizer is correctly set up too.

4.9 CF-Summary

This section is rather big, but at the same time it is not that many steps, much copy-paste. If everything goes well it should work immediately, if not, I recommend going through them one step at a time and trying to figure out if every value is set correct. In many of my copy-paste examples there are variables that needs to be replaced by your actual values. Be mindful!

- You validated YAML locally, uploaded children to S3, and created a single root stack that built the SG and EC2 layers in order.
- Outputs gave you public IPs; you derived private IPs for advertise-addr.
- You manually initialized/joined the swarm and deployed the same compose from [2.4](#).
- Result: identical outcome to the console walkthrough, but now fully reproducible via IaC + one root stack (create/update/delete).

When you are done, you have successfully converted the manual guide to Cloud Formation!

5. Introducing DynamoDB and reviewing IAM policies

Now we step out of the template-boundary. only real decision we did up until now, was to convert the manual setup to a Nested Stacks CloudFormation setup.

5.1 Design considerations

DynamoDB On-Demand (Pay per request) + strong perimeter has pros: zero capacity planning, autoscales up and down, but nasty cons: cost is proportional to traffic. A bot storm on our forms field could rack up charges! Cost-control and abuse protection is a must!

IMDSv2 – Instance MetaData Service v2 for EC2 is a local, hardened HTTP endpoint on each instance. That means each instance is attached to an IAM Role (Instance profile), with policies. Pros here is we need no hard-coded keys, credentials rotate automatically and every Swarm node inherits the same least-privilege access. Sounds neat?

IAM strategy – Least privilege is a must, and it has to function in a scalable environment. The idea is to attach one EC2 instance profile to all swarm nodes. IMDSv2 > STS (short lived credentials) > AWS SDK in the .NET app? No static keys is a win.

5.2 We have to begin somewhere. Begin the beginner

Plan: Attach an EC2 Role and enforce IMDSv2.

5.2.1 New child template: EC2 role + instance profile

Create infra/templates/20-iam-ec2-role.yaml

```
AWSTemplateFormatVersion: '2010-09-09'
Description: EC2 Instance Role + Instance Profile

Parameters:
  StackNamePrefix:
    Type: String

Resources:
  AppEc2Role:
    Type: AWS::IAM::Role
    Properties:
      RoleName: !Sub '${StackNamePrefix}-ec2-app-role'
      AssumeRolePolicyDocument:
        Version: '2012-10-17'
        Statement:
          - Effect: Allow
            Principal: { Service: ec2.amazonaws.com }
            Action: sts:AssumeRole

  AppEc2InstanceProfile:
    Type: AWS::IAM::InstanceProfile
    Properties:
      InstanceProfileName: !Sub '${StackNamePrefix}-ec2-app-profile'
      Roles: [ !Ref AppEc2Role ]

Outputs:
  InstanceProfileName:
```

```
Value: !Ref AppEc2InstanceProfile
```

Validate. Always [validate](#).

```
aws cloudformation validate-template --template-body file://infra/templates/20-iam-ec2-role.yaml
```

5.2.2 Attach the profile + enforce IMDSv2 on our existing EC2s

We need to modify our templates to attach the profile. Edit 10-ec2-swarm.yaml:

Under “Parameters” (probably line 4) you have six parameters. Add this after them:

```
InstanceProfileName:  
  Type: String
```

On each of the instances (under the Resources subsection, Manager should be on line 23, Worker1 and 2 on line 43 and 63 respectively), add these two properties:

```
IamInstanceProfile: !Ref InstanceProfileName  
MetadataOptions:  
  HttpEndpoint: enabled  
  HttpTokens: required # IMDSv2 required
```

It should look something like this:

```
16 |   Type: AWS::EC2::Image::Id  
17 |   InstanceProfileName:  
18 |     Type: String  
30 |  
31 |  
32 |  
33 |   SecurityGroupIds: [ !Ref SecurityGroupId ]  
34 |   IamInstanceProfile: !Ref InstanceProfileName  
35 |   MetadataOptions:  
     HttpEndpoint: enabled  
     HttpTokens: required # IMDSv2 required  
   Tags:
```

Guess what comes next? You’re right. Validate!

```
aws cloudformation validate-template --template-body file://infra/templates/10-ec2-swarm.yaml
```

5.2.3 Orchestrate this in root.yaml

First, we need to add the new 20-template to the nested stacks in the root template. We do this by opening root, scrolling down above the SwarmEc2Stack and add this code around line 45:

```
IamEc2RoleStack:  
  Type: AWS::CloudFormation::Stack  
  Properties:  
    TemplateURL: !Sub  
      'https://${AWS::Region}.amazonaws.com/${TemplateBucket}/${TemplatePrefix}20  
      -iam-ec2-role.yaml'  
  Parameters:  
    StackNamePrefix: !Ref StackNamePrefix
```

We need to pass the profile to the EC2 stack (the SwarmEc2Stack in root). To do so we have to update some parameters, notably we want to add “IamEc2RoleStack” in the DependsOn line, and then add a line at the end where we give the InstanceProfileName-parameter:

```
# Update this line
```

```

DependsOn: [ SwarmSecurityGroupStack, IamEc2RoleStack ]
    # Add this line to the end of SwarmEc2Stack
    InstanceProfileName: !GetAtt
IamEc2RoleStack.Outputs.InstanceProfileName

```

Validate. No really, just do it. You know you want to.

```
aws cloudformation validate-template --template-body file://infra/templates/root.yaml
```

5.2.4 Upload and update the root stack

This is a bit different from the earlier uploads and updates, now we have modified our EC2 Instances and attached a profile, so we run this first:

```

# Upload files to the S3 bucket
aws s3 cp infra/templates/20-iam-ec2-role.yaml "s3://${BUCKET}/${PREFIX}20-iam-ec2-role.yaml" --region
"${REGION}"
aws s3 cp infra/templates/10-ec2-swarm.yaml      "s3://${BUCKET}/${PREFIX}10-ec2-swarm.yaml"      --region
"${REGION}"
aws s3 cp infra/templates/root.yaml           "s3://${BUCKET}/${PREFIX}root.yaml"           --region
"${REGION}"

# Update root, instances will be replaced to attach the profile
aws cloudformation update-stack \
--region "${REGION}" --stack-name "${STACK}" \
--template-url "https://s3.${REGION}.amazonaws.com/${BUCKET}/${PREFIX}root.yaml" \
--parameters file://${PARAMS} \
--capabilities CAPABILITY_IAM CAPABILITY_NAMED_IAM

```

This might lead to a change of public IPs, because we replace the EC2s. Re-pull outputs:

```
aws cloudformation describe-stacks --region "${REGION}" --stack-name "${STACK}" \
--query 'Stacks[0].Outputs[].[OutputKey,OutputValue]' --output table
```

In my case I did get new IPs when doing this. I re-ran the [commands from 4.7.1](#) to store the new IPs in the variables.

I always keep a window open on AWS Console > CloudFormation > Stacks > stackname (swarm-cf-root in this case) so I can see live if the Swarms are healthy:

CloudFormation < **Stacks (4)**

Stack details

Drifts

StackSets

Exports

Infrastructure Composer

IaC generator

Hooks overview

Invocation summary **New**

Hooks

Registry

Public extensions

Activated extensions

Publisher

Stacks

NESTED
swarm-cf-root-lamEc2RoleStack-AQ2RUFXY365L
2025-09-12 12:00:06 UTC+0200
CREATE_COMPLETE

NESTED
swarm-cf-root-SwarmEc2Stack-1RSIQWJ8890CQ
2025-09-12 09:18:41 UTC+0200
UPDATE_COMPLETE

NESTED
swarm-cf-root-SwarmSecurityGroupStack-6TZK9YS4EKLK
2025-09-12 09:18:19 UTC+0200
UPDATE_COMPLETE

swarm-cf-root
2025-09-12 09:18:16 UTC+0200
UPDATE_COMPLETE

5.2.5 Verify IMDSv2 + role (on the Manager)

Let us start by SSHing into the Manager:

```
# SSH to the Manager
ssh -i ~/.ssh/tinfoil-eu-west-1.pem ec2-user@$MANAGER_PUB
# On the Manager, show the HTTP status
curl -o /dev/null -s -w "%{http_code}\n" http://169.254.169.254/latest/meta-data/iam/info
```

This should return a 401:

```
[ec2-user@ip-172-31-2-233 ~]$ curl -o /dev/null -s -w "%{http_code}\n" http://169.254.169.254/latest/meta-data/iam/info
401
```

401 means “unauthorized access”. We now need a token to access this information! Let us try get a token and then see if we can access the information:

```
# Grab a token
TOKEN=$(curl -sX PUT "http://169.254.169.254/latest/api/token" \
-H "X-aws-ec2-metadata-token-ttl-seconds: 21600")

# Use the token see if you can access information now
```

```
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/iam/security-credentials/
```

That is a solid response, mine returned `swarm-cf-ec2-app-role-etc!` Exit out of the Manager. You can now verify by checking Instance descriptions, they should say `HttpTokens: Required`. Try:

```
aws ec2 describe-instances --region "$REGION" --instance-ids "$MANAGER_ID" \
--query 'Reservations[0].Instances[0].MetadataOptions'
```

If you want you can also check if an instance profile is attached:

```
aws ec2 describe-instances --region "$REGION" --instance-ids "$MANAGER_ID" \
--query 'Reservations[0].Instances[0].IamInstanceProfile'
```

5.2.6 Summary and an explanation of what we just did

Wow, that is a lot of code spammage. We were not supposed to do that here, you promised you would explain what we are doing? Sorry, I was in a flow! Yes, you are right, so let us just do that. Normally I kept these summaries brief. This however, is a massive adaptation, so let us review:

What we added

- An IAM Role for EC2 + an Instance Profile (the “holder” that lets EC2 attach that role)
- We told all three swarm instances (manager + 2 workers) to use that profile (crucial!)
- We enforced IMDSv2 on the instances (`HttpTokens: required`)

How this works

- Think of the IAM Role as a keycard with permissions
- The Instance Profile is how that keycard is physically attached to each EC2 instance
- Inside each instance there is a tiny, local HTTP endpoint called IMDS (Instance Metadata Service)
- With IMDSv2, you must first ask for a short-lived token, then present that token to read metadata (like the role name) or to let the AWS SDK fetch temporary credentials for the role
- Your app (and AWS SDKs) automatically use IMDSv2 behind the scenes. You don’t code the token dance yourself—SDKs handle that. (We only used curl to prove it’s enforced)

Why we did it

- No hard-coded keys: the app will get short-lived credentials from the role. Nothing to store in code or `.env`
- Least privilege ready: the role currently has *no* data permissions. In the next step we’ll add just enough access to DynamoDB
- Secure by default: IMDSv2 blocks casual metadata access (and helps mitigate SSRF-style issues). Without the token, access is denied

What changed operationally

- Updating the template to attach the profile replaced the EC2 instances, so public IPs may have changed—we re-pulled outputs
- SSH still respects your `AllowedSshCidr`; if your VPN IP changes, update that parameter and `update-stack`

- From now on, any container scheduled by Swarm on any node has the same identity (the same role)—great for horizontal scaling

5.3 Let us introduce DynamoDB

Lorum Ipsum ad Infinitum

5.x Plan to be included:

Table for form submissions, S3 access patterns, instance/role permissions

Lorum Ipsum ad Infinitum

6. .NET, .NET, my Java-library for a .NET-app!

Plan to be included:

Minimal service that writes to S3/DynamoDB (behind the same IaC), then scale and front it later with ALB

Lorum Ipsum ad Infinitum

7. Where to go from here? (Upgrade the system)

Things we can consider:

- VPC Gateway Endpoints for DynamoDB (keep traffic in AWS backbone)
- ALB/ASG and a Bastion Host
- Private VPC, routing tables and multiple AZs
- Replacing DynamoDB with a proper DB or a storage contained in an Instance, should be more cost effective?
- Terraform instead of CloudFormation?
- SSM Session Manager? Then SSH port 22 can be removed?
- Obviously: CloudWatch, logging and metrics

Lorum Ipsum ad Infinitum

X. References

X.1 educ8.se

The basics of this guide is based on study material and tutorials handed out during the Cloud Developer YH-program (school of applied knowledge) at Campus Mölndal in 2025.

Copyright to that material belongs to our teacher Lars Appel, this guide has been modified (and I have taken personal design choices) to reflect this. Lars' guides are more handing out the basics, this one will try to explain why we chose what. See more at: <https://educ8me.se>

X.2 LLM support

I did use LLMs, not to write this text, it is my words, my explanations, and my architecture/solution with the nested stacks setup in CF.

However, when I converted the manual setup from section 2 to CloudFormation, I asked it to give me a template skeleton for CloudFormation to match the sections I wrote in part 2. I have built on that, but that saved me quite a lot of time of writing CF-templates.