

Index

1. Introduction	2
1.1 Pre-requisites	2
2. Infrastructure basics.....	3
2.1 Create a Security Group	3
2.1.1 Create the initial inbound rules	3
2.1.2 Update the Security Group	3
2.1.3 What the hell did we just do	4
2.1.4 Important note regarding SSH inbound:	4
2.1.5 Standard vs custom ports.....	4
2.1.6 TCP vs UDP protocols	5
2.1.7 SG-Summary.....	5
2.2 Launch EC2 Instances	5
2.2.1 First let us create the instances.....	5
2.2.1.1 ..and this script does what?	5
2.2.2 Summary, launch and multiples.....	6
2.2.3 Note Instance IPs.....	6
2.2.4 EC2-Summary.....	6
2.3 Initialize Docker Swarm via SSH	6
2.3.1 Create a nifty helper-document (optional but recommended)	6
2.3.2 Connect to the Manger and initialize Swarm.....	7
2.3.2.1 I got most of that. I think	7
2.3.3 Let us add the Worker nodes	7
2.3.3.1 Let us quickly verify the Swarm cluster.....	8
2.3.4 Swarm-Summary.....	8
2.4 Deploying the services manually	8
2.4.1 Create a Docker Compose file	8
2.4.1.1 Let us not go into all details, but roughly what this does is this	9
2.4.2 Deploy the freshly created Stack	9
2.4.3 Verify the deployment	10
2.4.4 Deploy-Summary.....	10
2.5 Test web service and see if the service can scale.....	10
2.5.1 Let us first verify web access.....	10
2.5.2 Scale the services	11

2.5.3 Tools to monitor the services.....	12
2.5.4 Web-Summary	12
2.6 Tools for cleaning up (optional).....	12
2.6.1 Remove stack	12
2.6.2 Leave Swarm	12
2.6.3 Terminate Instances	13
2.6.4 Delete Security Group.....	13
2.7 Summary of above sections	13
3 CloudFormation me up, Scotty	13
3.1 S3 Bucket roles first.....	14
3.1.1 Create the bucket first.....	14
3.1.2 Then attach the policy to the user so you can use the bucket.....	14
3.2 Quick conversion – CloudFormation style.....	15
3.2.1 Nested stacks structure.....	15
3.2.2 Upload child templates to S3	15
3.2.3 root.yaml (Master stack)	15
3.2.4 00-sg-swarm.yaml (Swarm Security Group).....	17
3.2.5 10-ec2-swarm.yaml (EC2 Manager + 2 Workers, Docker installed)	18
3.2.6 /parameters/dev.json (example).....	20
X. References.....	21
X.1 educ8.se	21

1. Introduction

This is just a tutorial of how we built a webform in .NET that talks through IAM/S3 to AWS DynamoDB and is using a Docker Swarm setup on a scalable AWS EC2-instance setup. I wrote this guide for my own good (learning by describing), but sharing it in case someone else might benefit from it too.

Link to repository: <https://github.com/mymh13/swarm-dotnet-test>

1.1 Pre-requisites

AWS account, AWS CLI installed, SSH key, AWS key pair, and basic cloud and .NET knowledge

2. Infrastructure basics

In this phase we will do the rudimentary basics: Create a security group for the swarm and launch EC2 instances. Further infrastructure (database) will be added later.

2.1 Create a Security Group

The SG sets the rules and boundaries for the system that uses that particular SG. In this first step, we want to create a SG that primarily handles the inbound rules for our Swarm network.

2.1.1 Create the initial inbound rules

In the AWS console, navigate to EC2 > Security Groups > Create Security Group. Fill in the following:

Name: Choose something descriptive so you know what role the SG plays. You will end up with many SGs eventually and then it will be hard to browse through them.

Description: Something descriptive, see above

VPC: use default VPC for this template, if you can do custom then you do not need this guide 😊

Inbound rules:

Type	Protocol	Port	Source	IP	Description
SSH	TCP	22	My IP	[yourIPhere]	SSH
HTTP	TCP	80	IPv4	0.0.0.0/0	HTTP inbound
Custom TCP	TCP	8080	IPv4	0.0.0.0/0	Visualizer

Create security group Info

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.

Basic details

Security group name Info
Tinfoil-Swarm-SG
Name cannot be edited after creation.

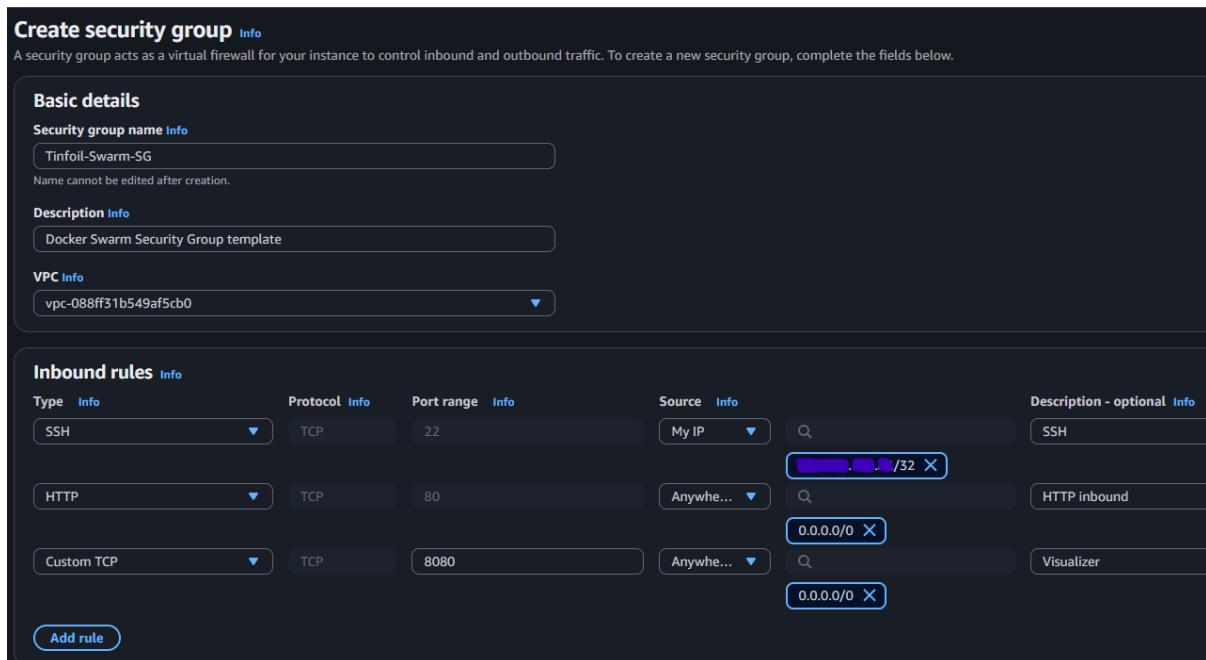
Description Info
Docker Swarm Security Group template

VPC info
vpc-088ff31b549af5cb0

Inbound rules Info

Type	Protocol	Port range	Source	Description - optional
SSH	TCP	22	My IP	SSH
HTTP	TCP	80	Anywhere...	HTTP inbound
Custom TCP	TCP	8080	Anywhere...	Visualizer

Add rule



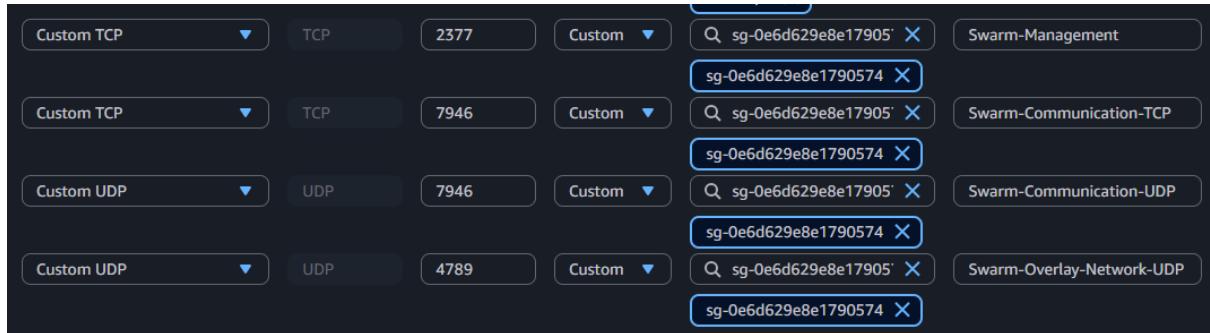
At this point you want to create the security group, there will be four more inbound rules but they will use self-reference so you want to create the SG before it can refer to itself!

2.1.2 Update the Security Group

We need to add the final inbound rules. Edit the Security Group. In the IP/search window, click the window and choice your SG you just created to self-reference.

Inbound rules (3)							
	Name	Security group rule ID	IP version	Type	Protocol	Port range	Source
<input type="checkbox"/>	-	sgr-0e5bb9118013acf06	IPv4	Custom TCP	TCP	8080	0.0.0.0/0 Visualizer
<input type="checkbox"/>	-	sgr-0261101699f05fb43	IPv4	SSH	TCP	22	185.209.198.82/32 SSH
<input type="checkbox"/>	-	sgr-075ac86c81497598d	IPv4	HTTP	TCP	80	0.0.0.0/0 HTTP inbound

Do note UDP on the two last ones! It should look something like this:



2.1.3 What the hell did we just do

Before we move on, let us briefly explain what this does:

- SSH (TCP 22) – Lets you connect to the server via SSH. We lock it down to your own IP only, so no one else can try to log in.
- HTTP (TCP 80) – This allows anyone on the internet to reach our application in the browser.
- Visualizer (TCP 8080) – Docker Swarms visualizer tool runs on port 8080. This rule makes us able to reach and view the visualizer in our browser.
- Swarm Management (TCP 2377) – Used for communication between the swarm manager and the worker nodes. Self-referenced so only other servers in the same SG are allowed in.
- Swarm Communication (TCP 7946 + UDP 7946) – These ports handle internal gossip (yes, that is the actual term in distributed systems! It means each node spread information to its neighbours) and node discovery within the system. Self-reference = self-contained, only swarm members allowed.
- Overlay Network (UDP 4789) – This port carries actual container traffic between swarm nodes over the same overlay network. Self-referenced, internal traffic inside the swarm.

2.1.4 Important note regarding SSH inbound:

Personally, I am always on VPN and my IP changes frequently. This means I have to consider how to handle my SSH IP inbound. In a brief project like this I chose to edit the SSH rule and adjust to reflect my new IP now and then, but for a longer term – and especially if you are more than one person handling this node – you need a more robust solution.

2.1.5 Standard vs custom ports

- Standard: TCP 22 is standard SSH-port. TCP 80 is standard HTTP- (web) traffic.
- Widely used: TCP 8080 is not an official standard, but a widely used convention for dashboards and dev tools.
- Custom: TCP 2377, TCP 7946 + UDP 7946 and UDP 4789 are specific to Docker Swarm.

Custom ports are chosen by the system- or application designer, it is important to know the standards and commonly used ones, so we don't accidentally re-use them for custom roles.

2.1.6 TCP vs UDP protocols

Most internet traffic uses one of the two protocols: TCP or UDP.

- TCP (Transmission Control Protocol) – It sets up a connection, checks that messages arrive, re-sends everything that gets lost. It is reliable but has a bit more overhead. We use this for SSH and HTTP because accuracy is more important than speed.
- UDP (User Datagram Protocol) – It is sent without checking if it arrives, thus making it much faster and uses less overhead than TCP, but delivery is not guaranteed! This is useful for Docker Swarm's internal networking, especially if we are working idempotent as we can accept lost packets.

In short: TCP for reliable traffic, UDP for fast cluster chatter and data transport.

2.1.7 SG-Summary

We have created a Security Group that will handle inbound rules, boundaries and internal traffic.

2.2 Launch EC2 Instances

EC2 Instances is basically just Virtual Machines (VM) but let us stick to AWS-terminology. They will work as “servers” (nodes/hosts might be more appropriate description) for our Swarm. At this point we could also have considered other VMs for roles like a Bastion Host, or a Database or something else, but let us keep this tutorial simple for now and just focus on the Docker Swarm.

2.2.1 First let us create the instances

AWS Console > EC2 > Instances > Launch Instance

Name: Just like with the SG group and classes/methods, always use descriptive names!

AMI: Amazon Linux 2023

Instance type: t3.small

Key pair: This was a pre-req 😊 chose your already pre-set-AWS-key-pair for this

Network settings:

- VPC – default VPC
- Subnet – default Subnet
- Auto-assign public IP – enable
- Firewall / Security Group – Select existing (the one you just created in step 1)

Advanced details > User Data: add a start script for the instance:

```
#!/bin/bash
dnf update -y
dnf install -y docker
systemctl enable --now docker
usermod -aG docker ec2-user
```

2.2.1.1 ..and this script does what?

`dnf update -y` - Updates all existing software on the instance. The `-y` flag means “say yes to all prompts” so the process doesn't stop to ask

`dnf install -y docker` - Installs Docker from the Amazon Linux package repository.
Again, -y auto-confirms

`systemctl enable --now docker` - Tells the system to start Docker right away (--now) and to also start it automatically on every reboot (enable)

`usermod -aG docker ec2-user` - Adds the default AWS login user (ec2-user) to the docker group, so you can run Docker commands without needing sudo every time

2.2.2 Summary, launch and multiples

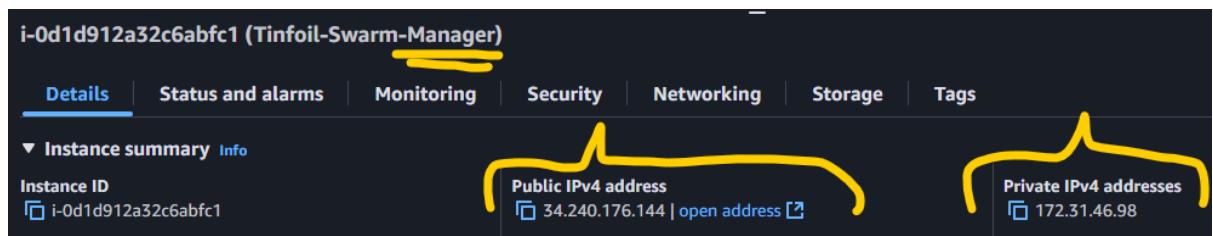
On your right hand, you should have a “Summary” window, with a box named “number of instances” – chose 3 and then click Launch instance.

Then go your EC2 > Instances and rename the second and third Manager-instances to Worker-1 + 2:

Instances (3) Info							Last updated 2 minutes ago	Connect	Instance state ▾	Actions ▾	Launch instances
<input type="checkbox"/> Name ⚙		Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone				
<input type="checkbox"/>	Tinfoil-Swarm-Manager	i-0d1d912a32c6abfc1	Running Q Q	t3.small	Initializing	View alarms +	eu-west-1a				
<input type="checkbox"/>	Tinfoil-Swarm-Worker-1	i-03bcaf2d6f7c8f0dc	Running Q Q	t3.small	Initializing	View alarms +	eu-west-1a				
<input type="checkbox"/>	Tinfoil-Swarm-Worker-2	i-02c3341722808b0fd	Running Q Q	t3.small	Initializing	View alarms +	eu-west-1a				

2.2.3 Note Instance IPs

View the instances (in the version above, September 2025, you can just click the instance tick box to the left and you get a quick view of the instance in a window below the instances). You want to make sure you have the Public IP and the Private IP for the three instances. I copy-pasted them down in a markdown document for easy access.



2.2.4 EC2-Summary

We have set up three EC2 Instances, one working as a Master and two Workers nodes.

2.3 Initialize Docker Swarm via SSH

I hope you have got the IPs handy and SSH + Key Pair is set up. In this phase we will initialize Docker Swarm on our EC2 Instances and verify the cluster. Let's go.

2.3.1 Create a nifty helper-document (optional but recommended)

I used Obsidian and wrote in markdown, but anything will do. What I did was having a number of sections where I listed the public/private IPs of the Instances and a copy-paste command ready to SSH into the Instances. It should look something like this:

Swarm Manager

```
- public IP - 85.482.592.492  
- private IP - 172.31.46.98  
ssh -i ~/.ssh/<your-key-name>.pem ec2-user@85.482.592.492
```

Do note:

- This assumes your SSH key is in the default user/.ssh/ directory. If you are a Windows user I can highly recommend to always SSH keys in the C:\Users\<username>\.ssh directory
- The public IP in the example above has randomized numbers, this is information you should protect so make sure you do not save this information open in – for example – your project folder (if you use that for a GitHub repo or similar)

2.3.2 Connect to the Manager and initialize Swarm

Use the SSH command [found in 2.3.1](#) to SSH into your Manager node:

```
ssh -i ~/.ssh/your-key.pem ec2-user@<manager-public-ip>
```

Then we will initialize the Swarm on the Manager:

```
sudo docker swarm init --advertise-addr <manager-private-ip>
```

We should get a verification that looks similar to this:

```
Swarm initialized: current node (xyz123) is now a manager.
```

To add a worker to this swarm, run the following command:

```
docker swarm join --token SWMTKN-1-xxx... <manager-private-ip>:2377
```

Copy the join-command and save it to your helper-document.

```
sudo docker swarm join --token SWMTKN-1-<long-string> <manager-private-ip>:2377
```

2.3.2.1 I got most of that. I think

No worries, let me elaborate what you just did.

- By SSHing into the Manager node and running docker swarm init, we turned that server into the “brain” of the swarm.
- The --advertise-addr <manager-private-ip> flag makes sure other servers connect over the private AWS network (faster and more secure than public IP).
- Docker then gave us a join token and command. This is like a one-time password that worker nodes will use to join the swarm. Copy it somewhere safe — you’ll need it in the next step.

2.3.3 Let us add the Worker nodes

Basically, we will join the existing Swarm as worker nodes, there is not much else to it. Open a new terminal window and run these commands:

```
ssh -i ~/.ssh/your-key.pem ec2-user@<worker1-public-ip>  
# Run the join command from step 2.3.2  
sudo docker swarm join --token SWMTKN-1-xxx... <manager-private-ip>:2377
```

Repeat the above step but replace worker1-public-ip with worker2's public IP. I do not think we need to explain the command in-depth, as you see we use "swarm join" which should be self-explanatory.

You should have three terminal windows that look something like this:

```
~/
```

```
[ec2-user@ip-172-31-46-98 ~]$ sudo docker swarm init --advertise-addr 172.31.46.98
Swarm initialized: current node (sfzku1d14tps762bvhufz89t) is now a manager.

To add a worker to this swarm, run the following command:

  docker swarm join --token SWMTKN-1-32we17ks1h0hkee5q3t3virt57d74cjieti9kfivtksfq1nenn-7
  72.31.46.98:2377
This node joined a swarm as a worker.
[ec2-user@ip-172-31-39-19 ~]$ ...
~/
```

```
[ec2-user@ip-172-31-42-153 ~]$ docker swarm join --token SWMTKN-1-32we17ks1h0hkee5q3t3virt57d74cjieti9kfivtksfq1nenn-7
  72.31.46.98:2377
This node joined a swarm as a worker.
[ec2-user@ip-172-31-42-153 ~]$ ...
```

2.3.3.1 Let us quickly verify the Swarm cluster

In the Manager nodes' terminal, type this: `sudo docker node ls`

You should see this:

```
[ec2-user@ip-172-31-46-98 ~]$ sudo docker node ls
ID           HOSTNAME   STATUS  AVAILABILITY  MANAGER STATUS  ENGINE VERSION
7dyzmrcn3ausuqq0mjf8we65f  ip-172-31-39-19.eu-west-1.compute.internal  Ready  Active
1eahizoejo58fet83gfmkgo6e  ip-172-31-42-153.eu-west-1.compute.internal  Ready  Active
sfzku1d14tps762bvhufz89t *  ip-172-31-46-98.eu-west-1.compute.internal  Ready  Active      Leader
25.0.8
```

2.3.4 Swarm-Summary

We have initialized a Docker Swarm on our Manager, joined with Workers and verified it.

2.4 Deploying the services manually

Ideally this is something that we have an automatic solution for, some IaC solution would be preferable. But this is just a basic template to grasp the concept so we will deploy manually.

2.4.1 Create a Docker Compose file

On the Manager node, we will create the stack file by typing (pasting) this:

```
cat > docker-stack.yml << 'EOF'
version: "3.8"
services:
  web:
    image: nginx:stable-alpine
    deploy:
```

```

replicas: 3
restart_policy:
  condition: on-failure
update_config:
  parallelism: 1
  delay: 5s
ports:
  - "80:80"
networks: [webnet]

viz:
  image: dockersamples/visualizer:stable
deploy:
  placement:
    constraints: [node.role == manager]
  ports:
    - "8080:8080"
  volumes:
    - /var/run/docker.sock:/var/run/docker.sock
  networks: [webnet]

networks:
  webnet:
    driver: overlay
EOF

```

2.4.1.1 Let us not go into all details, but roughly what this does is this

The cat > docker-stack.yml << 'EOF' ... EOF command is a way to create a new file (docker-stack.yml) and paste content directly into it. Everything between the two EOF markers becomes the file content.

This file is written in YAML, which is very sensitive to spacing and indentation — so it's best to copy-paste rather than type it manually.

The stack defines two services:

- Web: An nginx server with 3 replicas, exposed on port 80
- Viz: A Docker Swarm visualizer, only running on the manager node, exposed on port 8080 (we set that port in the SG but the Manager runs the Viz itself)

At the bottom, it also defines a custom overlay network (webnet) that lets the services talk to each other across swarm nodes.

2.4.2 Deploy the freshly created Stack

Still on the Manager node, run the deploy command:

```
sudo docker stack deploy -c docker-stack.yml myappname
```

```
[ec2-user@ip-172-31-46-98 ~]$ sudo docker stack deploy -c docker-stack.yml myappname
Creating network myappname_webnet
Creating service myappname_viz
Creating service myappname_web
```

This tells Docker Swarm to deploy the stack file (-c docker-stack.yml) and give the stack the name myappname.

2.4.3 Verify the deployment

```
# Show all stacks currently running in the Swarm
sudo docker stack ls

# List the services inside the stack, along with how many replicas are
running. If we followed the guide you should see web with 3/3 replicas and viz
with 1/1 - since we have 3x instances but only the Manager runs the viz.
sudo docker service ls

# Detailed info about each service: which node it is running on, current
state, specific container ID. This is useful when troubleshooting.
sudo docker service ps myappname_web
sudo docker service ps myappname_viz
```

The comment above each command explains them further in-depth.

```
[ec2-user@ip-172-31-46-98 ~]$ sudo docker stack ls
NAME      SERVICES
myappname  2

[ec2-user@ip-172-31-46-98 ~]$ sudo docker service ls
ID          NAME      MODE      REPLICAS      IMAGE                                     PORTS
orr1o0zokn16  myappname_viz  replicated  1/1  dockersamples/visualizer:stable  *:8080->8080/tcp
oq6uit43aetn  myappname_web  replicated  3/3  nginx:stable-alpine           *:80->80/tcp

[ec2-user@ip-172-31-46-98 ~]$ sudo docker service ps myappname_web
ID          NAME      IMAGE      NODE      DESIRED STATE  CURRENT STATE      ERROR      PORTS
pvilkvk75478  myappname_web.1  nginx:stable-alpine  ip-172-31-46-98.eu-west-1.compute.internal  Running   Running  3 minutes ago
z78t0hc6xyfj  myappname_web.2  nginx:stable-alpine  ip-172-31-42-153.eu-west-1.compute.internal  Running   Running  3 minutes ago
eg510zadhb03  myappname_web.3  nginx:stable-alpine  ip-172-31-39-19.eu-west-1.compute.internal  Running   Running  3 minutes ago
[ec2-user@ip-172-31-46-98 ~]$ sudo docker service ps myappname_viz
ID          NAME      IMAGE      NODE      DESIRED STATE  CURRENT STATE      ERROR      PORTS
nzat3lxmepib  myappname_viz.1  dockersamples/visualizer:stable  ip-172-31-46-98.eu-west-1.compute.internal  Running   Running  3 minutes ago
```

2.4.4 Deploy-Summary

We have created and deployed the Docker Compose-file on the Manager and verified it.

2.5 Test web service and see if the service can scale

Now let's test that the web service is accessible, and then try scaling it up and down.

2.5.1 Let us first verify web access

Open a browser and go to the public IP of any node:

<http://<public-ip-of-any-node>>

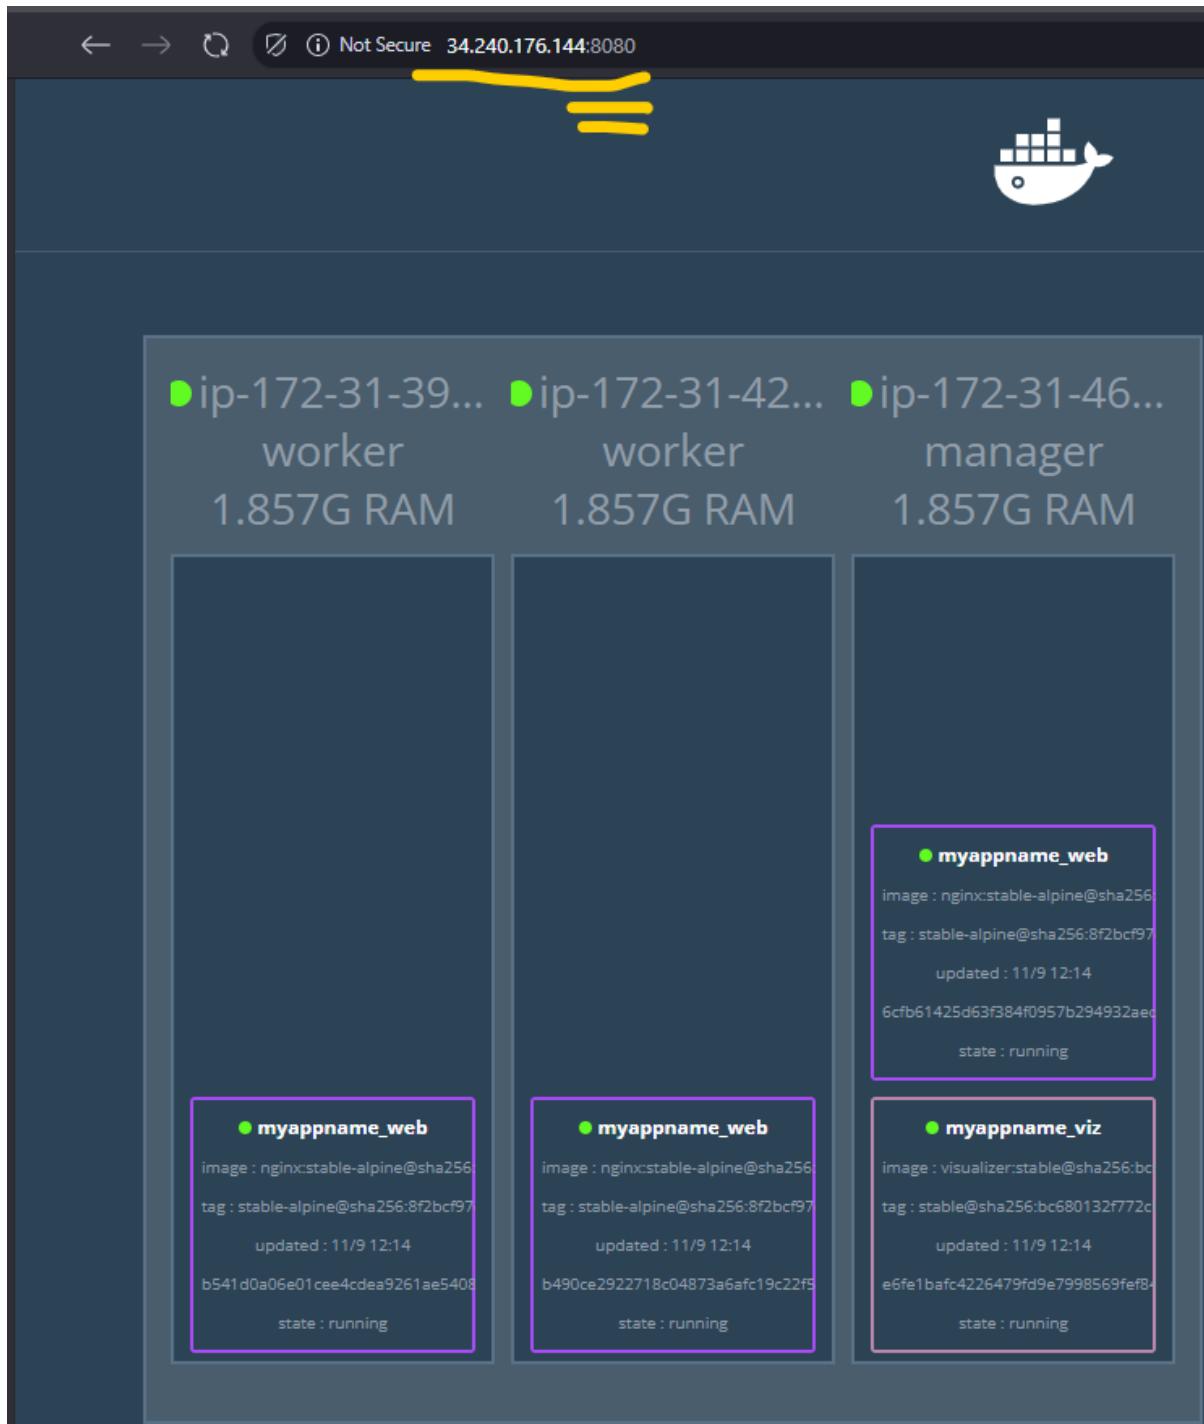
You should see the default Nginx welcome page. Then check the visualizer service:

<http://<manager-public-ip>:8080/>

If everything is configured correctly, you'll see a live visualization of your swarm.

- Port 80 needs to be open for web access.
- Port 8080 needs to be open for the visualizer.

If either page doesn't load, double-check your Security Group and your stack file to make sure the correct ports are open to the right IPs.



2.5.2 Scale the services

Keep that Visualizer browser tab open, and run these commands on the Manager node:

```
# Scale up nginx to 5 replicas
sudo docker service scale myappname_web=5

# Check scaling
sudo docker service ps myappname_web

# Scale back down to 3
sudo docker service scale myappname_web=3
```

You should see containers appear and disappear in the visualizer as the service scales up and down.

Do note: Scaling here means adding or removing containers, not EC2 instances. Your swarm is still running on the same set of servers. You're just changing how many copies of the nginx container run across them.

2.5.3 Tools to monitor the services

You won't need these right now, as you saw the services running, but leaving this here as it is helpful tools in the future:

```
# Watch service status (Ctrl + C to exit)
watch sudo docker service ls

# View service logs
sudo docker service logs myappname_web
sudo docker service logs myappname_viz
```

"Watch" is live monitoring, so you will have to Ctrl + C to exit / cancel the operation.

Logs are useful tools, as a developer these are your best friends. Logging separate logfiles for separate services / functions is something you will want to do and learn.

2.5.4 Web-Summary

We visit the public IPs, check that the service is running and try the Visualizer.

2.6 Tools for cleaning up (optional)

At this point you might want to clean up a stack (if you are running one), or leave / reset the Swarm.

2.6.1 Remove stack

On the manager node: `sudo docker stack rm myappname`

`rm` means "remove", it is a good command to know if you are working with files and directories too.

2.6.2 Leave Swarm

You might want to reset the swarm, alter instances or something else:

```
# On workers
sudo docker swarm leave
```

```
# On manager (force)
sudo docker swarm leave --force
```

2.6.3 Terminate Instances

Visit EC2 > Instances. Select all three instances. Under “Instance state” chose Terminate.

Name	Instance ID	Instance state	Instance type
Tinfoil-Swarm-Manager	i-0d1d912a32c6abfc1	Running	t3.small
Tinfoil-Swarm-Worker-1	i-03bcf2d6f7c8f0dc	Running	t3.small
Tinfoil-Swarm-Worker-2	i-02c3341722808b0fd	Running	t3.small

2.6.4 Delete Security Group

Visit EC2 > Security Groups. Under Actions, Delete the SG.

Name	Security group ID	Description
-	sg-022ed4a7f9cbcf8d1	default
<input checked="" type="checkbox"/> -	sg-0e6d629e8e1790574	Tinfoil-Swarm-SG

2.7 Summary of above sections

We have built a SG, launched three instances and deployed a Docker Swarm on them, making sure it is healthy, and all the roles are functional.

In the class when we were at this point, we were supposed to do a .NET application and use DynamoDB for handling files that we send (and possibly retrieve, but that was a bonus) from our .NET application. As I have taken down and up these swarms now multiple times, I feel it is appropriate to create section 3 here – create the above system but in CloudFormation (or Terraform, but let us begin with CF).

3 CloudFormation me up, Scotty

I will build this using the Nested Stacks system as I really, strongly dislike messy long files. It has other benefits too:

- Separations of Concern – Each section is independent from other sections. This is good both from a modularity concept but it also easier when trying to solve problems that arise
- Deploy – Deployment benefits greatly from being separated. You can work on one module and just deploy that one, especially if we are working idempotent
- Adaptability – We might want to expand, or reduce, the program. Adapt to various needs. Then it might be beneficial to just replace or remove or add sections that we need.

That said, let us get going!

3.1 S3 Bucket roles first

Before CloudFormation can deploy nested stacks, we need a storage location for the templates and a user that has permissions to interact with that location. This is done in two parts.

3.1.1 Create the bucket first

1. Go to AWS Console > S3 > Create bucket
2. Enter a unique bucket name (ex., cf-artifacts-<account-id>-eu-west-1). Bucket names must be globally unique
3. Select the region where you plan to deploy your stacks (ex., eu-west-1)
4. Leave other settings at defaults for now
5. Click Create bucket

You now have a dedicated bucket for storing CloudFormation templates and artifacts.

3.1.2 Then attach the policy to the user so you can use the bucket

1. Go to AWS Console > IAM > Users > select your user
2. Click Add Permissions > Create inline policy
3. Switch to the JSON tab and paste in your policy (replace <your-bucket-name> with the one you just created in 3.1.1.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3>ListBucket",
        "s3:GetBucketLocation"
      ],
      "Resource": "arn:aws:s3:::<your-bucket-name>"
    },
    {
      "Effect": "Allow",
      "Action": [
        "s3GetObject",
        "s3PutObject",
        "s3DeleteObject",
        "s3DeleteObjectVersion"
      ],
      "Resource": "arn:aws:s3:::<your-bucket-name>/*"
    }
  ]
}
```

```
        }
    ]
}
```

Name the policy something descriptive (ex. S3_bucket_policy or cf-s3access-policy) and save it. Your IAM user now has the necessary permissions to interact with the S3 bucket – congrats!

3.2 Quick conversion – CloudFormation style

Before we write anything further: this document has the potential to become rather long from here on. So I will handle it this way – the skeleton of a nested stacks CloudFormation for the above project (section 1 + 2) will be posted here in plain text. But those documents will be growing, and we will add features.

That will too long to handle in text, so I recommend visiting the repo for the latest code:

<https://github.com/mymh13/swarm-dotnet-test>

3.2.1 Nested stacks structure

Child templates live in S3 and a tiny **root** template orchestrates them in order.

```
/infra
  /templates
    root.yaml          # orchestrates children
    00-sg-swarm.yaml   # Security Group for swarm
    10-ec2-swarm.yaml  # EC2 Manager + 2 Workers (Docker installed)
  /parameters
    dev.json           # example parameter set for root.yaml
```

3.2.2 Upload child templates to S3

We upload the SG and EC2-template stacks to our S3-bucket. root has defined URLs to the template stacks. When we run the aws update-stack-command on root it will update with the latest files that we have in the bucket. This way we can do manual work on separate stacks, but the root will orchestrate. This also means we can deploy and take down the entire stack with the root.

3.2.3 root.yaml (Master stack)

This will be a long “skeleton” code, and this all assumes you have basic understanding of AWS and Cloudformation. The SG/EC2-setup we run should be self-explanatory more or less though. Template:

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Root stack for Docker Swarm (manager + 2 workers) using nested stacks

Parameters:
  StackNamePrefix:
    Type: String
    Default: swarm-cf
    Description: Prefix used for naming
  TemplateBucket:
```

```

Type: String
Description: S3 bucket that stores child templates (e.g. cf-artifacts-<acct>-<region>)
TemplatePrefix:
  Type: String
  Default: swarm-iac/templates/
  Description: Key prefix inside the S3 bucket
VpcId:
  Type: AWS::EC2::VPC::Id
  Description: VPC to deploy into (choose your default VPC)
PublicSubnetId:
  Type: AWS::EC2::Subnet::Id
  Description: Public subnet for all three instances (default subnet is fine)
AllowedSshCidr:
  Type: String
  Description: Your IP in CIDR for SSH (e.g. 1.2.3.4/32)
KeyPairName:
  Type: AWS::EC2::KeyPair::KeyName
  Description: Existing EC2 key pair name
InstanceType:
  Type: String
  Default: t3.small
AmiId:
  Type: AWS::EC2::Image::Id
  Description: Amazon Linux 2023 AMI in this region

Resources:
SwarmSecurityGroupStack:
  Type: AWS::CloudFormation::Stack
  Properties:
    TemplateURL: !Sub
      'https://s3.${AWS::Region}.amazonaws.com/${TemplateBucket}/${TemplatePrefix}00-sg-swarm.yaml'
    Parameters:
      StackNamePrefix: !Ref StackNamePrefix
      VpcId: !Ref VpcId
      AllowedSshCidr: !Ref AllowedSshCidr

SwarmEc2Stack:
  Type: AWS::CloudFormation::Stack
  DependsOn: SwarmSecurityGroupStack
  Properties:
    TemplateURL: !Sub
      'https://s3.${AWS::Region}.amazonaws.com/${TemplateBucket}/${TemplatePrefix}10-ec2-swarm.yaml'
    Parameters:
      StackNamePrefix: !Ref StackNamePrefix

```

```

    PublicSubnetId: !Ref PublicSubnetId
    SecurityGroupId: !GetAtt
SwarmSecurityGroupStack.Outputs.SwarmSecurityGroupId
        KeyPairName: !Ref KeyPairName
        InstanceType: !Ref InstanceType
        AmiId: !Ref AmiId

Outputs:
    ManagerPublicIp:
        Value: !GetAtt SwarmEc2Stack.Outputs.ManagerPublicIp
    Worker1PublicIp:
        Value: !GetAtt SwarmEc2Stack.Outputs.Worker1PublicIp
    Worker2PublicIp:
        Value: !GetAtt SwarmEc2Stack.Outputs.Worker2PublicIp
    SecurityGroupId:
        Value: !GetAtt SwarmSecurityGroupStack.Outputs.SwarmSecurityGroupId

```

3.2.4 00-sg-swarm.yaml (Swarm Security Group)

This stack controls the Swarms SG.

```

AWSTemplateFormatVersion: '2010-09-09'
Description: Security Group for Swarm (SSH, HTTP, Viz, Swarm ports)

Parameters:
    StackNamePrefix:
        Type: String
    VpcId:
        Type: AWS::EC2::VPC::Id
    AllowedSshCidr:
        Type: String

Resources:
    SwarmSG:
        Type: AWS::EC2::SecurityGroup
        Properties:
            GroupDescription: !Sub '${StackNamePrefix}-swarm-sg'
            VpcId: !Ref VpcId
            SecurityGroupIngress:
                # SSH (locked to your IP)
                - IpProtocol: tcp
                  FromPort: 22
                  ToPort: 22
                  CidrIp: !Ref AllowedSshCidr
                # HTTP 80 (world)
                - IpProtocol: tcp
                  FromPort: 80
                  ToPort: 80

```

```

        CidrIp: 0.0.0.0/0
    # Visualizer 8080 (world)
    - IpProtocol: tcp
      FromPort: 8080
      ToPort: 8080
      CidrIp: 0.0.0.0/0
    # Swarm mgmt 2377 (self-reference)
    - IpProtocol: tcp
      FromPort: 2377
      ToPort: 2377
      SourceSecurityGroupId: !Ref SwarmSG
    # Gossip TCP 7946 (self-reference)
    - IpProtocol: tcp
      FromPort: 7946
      ToPort: 7946
      SourceSecurityGroupId: !Ref SwarmSG
    # Gossip UDP 7946 (self-reference)
    - IpProtocol: udp
      FromPort: 7946
      ToPort: 7946
      SourceSecurityGroupId: !Ref SwarmSG
    # Overlay UDP 4789 (self-reference)
    - IpProtocol: udp
      FromPort: 4789
      ToPort: 4789
      SourceSecurityGroupId: !Ref SwarmSG
Tags:
- Key: Name
  Value: !Sub '${StackNamePrefix}-swarm-sg'

Outputs:
SwarmSecurityGroupId:
  Value: !Ref SwarmSG

```

3.2.5 10-ec2-swarm.yaml (EC2 Manager + 2 Workers, Docker installed)

Here we create the EC2-Swarm, and set the User Data to install docker, like we did [here](#).

```

AWSTemplateFormatVersion: '2010-09-09'
Description: EC2 instances for Swarm (1 manager, 2 workers) with Docker
installed

Parameters:
StackNamePrefix:
  Type: String
PublicSubnetId:
  Type: AWS::EC2::Subnet::Id
SecurityGroupId:

```

```

Type: String
KeyPairName:
  Type: AWS::EC2::KeyPair::KeyName
InstanceType:
  Type: String
AmiId:
  Type: AWS::EC2::Image::Id

Mappings: {}

Resources:
  Manager:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: !Ref AmiId
      InstanceType: !Ref InstanceType
      KeyName: !Ref KeyPairName
      SubnetId: !Ref PublicSubnetId
      SecurityGroupIds: [ !Ref SecurityGroupId ]
    Tags:
      - Key: Name
        Value: !Sub '${StackNamePrefix}-manager'
  UserData:
    Fn::Base64: !Sub |
      #cloud-config
      runcmd:
        - dnf update -y
        - dnf install -y docker
        - systemctl enable --now docker
        - usermod -aG docker ec2-user

Worker1:
  Type: AWS::EC2::Instance
  Properties:
    ImageId: !Ref AmiId
    InstanceType: !Ref InstanceType
    KeyName: !Ref KeyPairName
    SubnetId: !Ref PublicSubnetId
    SecurityGroupIds: [ !Ref SecurityGroupId ]
  Tags:
    - Key: Name
      Value: !Sub '${StackNamePrefix}-worker-1'
  UserData:
    Fn::Base64: !Sub |
      #cloud-config
      runcmd:
        - dnf update -y
        - dnf install -y docker

```

```

        - systemctl enable --now docker
        - usermod -aG docker ec2-user

Worker2:
Type: AWS::EC2::Instance
Properties:
  ImageId: !Ref AmiId
  InstanceType: !Ref InstanceType
  KeyName: !Ref KeyPairName
  SubnetId: !Ref PublicSubnetId
  SecurityGroupIds: [ !Ref SecurityGroupId ]
Tags:
  - Key: Name
    Value: !Sub '${StackNamePrefix}-worker-2'
UserData:
  Fn::Base64: !Sub |
    #cloud-config
    runcmd:
      - dnf update -y
      - dnf install -y docker
      - systemctl enable --now docker
      - usermod -aG docker ec2-user

Outputs:
ManagerPublicIp:
  Value: !GetAtt Manager.PublicIp
Worker1PublicIp:
  Value: !GetAtt Worker1.PublicIp
Worker2PublicIp:
  Value: !GetAtt Worker2.PublicIp

```

Note:

This intentionally stops at Docker installed. After stack completes, follow the tutorial's SSH steps:

- SSH to manager, run docker swarm init --advertise-addr <manager-private-ip>.
- SSH to each worker, run the docker swarm join ... command from the manager output.
- Deploy your docker-stack.yml exactly as before.

3.2.6 /parameters/dev.json (example)

This is just a brief section that introduces the dev.json parameters.

X. References

X.1 educ8.se

The basics of this guide is based on study material and tutorials handed out during the Cloud Developer YH-program (school of applied knowledge) at Campus Mölndal in 2025. Copyright to that material belongs to our teacher Lars Appel, this guide has been modified (and I have taken personal design choices) to reflect this. Lars' guides are more handing out the basics, this one will try to explain why we chose what. See more at: <https://educ8me.se>

X.2